A linked resource of microbial genomes, community profiles and environmental data from the Rub' al-Khali

Rund Tawfiq[1,9], Marwa Abdelhakim[1], Sulaiman M. Alajel[2], Mohammed Alarawi[1], Hind Aldakhil[3], Abderahmane Derouiche[1], Daniela I. Drautz-Moses[4], Michel Dumontier[5], Raik Grünberg[6], Maxat Kulmanov[1], Alejandra Lopez-Velazquez[1], Susana Martinez Arbas[7], Kexin Niu[1], Krishnakumar Sivakumar[1], Tiannyu Wang[8], Xiang Zhao[4], Jood Zubair[1], Magnus Rueping[8], Robert Hoehndorf[1*]

[1]Bio-Ontology Research Group (BORG), Computer, Electrical and Mathematical Sciences and Engineering (CEMSE) Division, King Abdullah University of Science and Technology (KAUST), Thuwal, Saudi Arabia.
[2]Ministry of Environment, Water and Agriculture, Saudi Arabia.
[3]National Livestock & Fisheries Development Program (NLFDP), Ministry of Environment, Water and Agriculture (MEWA), Riyadh, 13413, Saudi Arabia.
[4]Bioscience Core Lab, King Abdullah University of Science and Technology, Thuwal, Saudi Arabia.
[5]Institute of Data Science, Department of Advanced Computing Sciences, Faculty of Science and Engineering, Maastricht University, Maastricht, The Netherlands.
[6]Biomedical Sciences Division (BioMed), King Abdullah University of Science and Technology (KAUST), Thuwal, Saudi Arabia.
[7]Luxembourg Centre for Systems Biomedicine (LCSB), University of Luxembourg, Esch-sur-Alzette, Luxembourg.
[8]Physical Science and Engineering (PSE) Division, King Abdullah University of Science and Technology (KAUST), Thuwal, Saudi Arabia.
[9]Sano Centre for Computational Personalized Medicine, Kraków, Poland.

*Corresponding author(s). E-mail(s): robert.hoehndorf@kaust.edu.sa;

Abstract

Repeated sampling across a 1,043 km transect of the Saudi Rub' al-Khali produced a linked resource of microbial community profiles, genomes and environmental observations. Five campaigns contributed 2,516 non-control specimen entries from 2,550 source records. The resource connects these specimens to laboratory processing, a canonical amplicon table of 1,271 profiles (1,237 after quality filtering), 551 shotgun libraries and 5,438 species-representative metagenome-assembled genomes. Genome products include quality estimates, taxonomy, gene annotations and per-library coverage. Field observations, archived-soil pH, X-ray fluorescence measurements and climate summaries retain their acquisition methods and quality dispositions. A modular RDF knowledge base uses shared identifiers to preserve specimen identity and assay-specific provenance across the records. Validation checks source reconciliation, measurement admission, identifier mappings, structural constraints and query outputs. Versioned graph modules, tabular and sequence products, and workflow code support reuse through archived files, a web portal and a SPARQL endpoint. The accompanying inventories distinguish the asserted graph from associated biological products and analysis-specific cohorts.

Keywords: Knowledge graph, Amplicon sequencing, Shotgun metagenomics, Metagenome-assembled genomes, Rub' al Khali, SIO, Semantic web

1 Background and summary

The Rub' al-Khali (Empty Quarter) is a hyper-arid sand desert that spans approximately 500,000 km^2 of the southern Arabian Peninsula[1]. Its landscape is dominated by extensive linear dunes known as *'uruq* (singular *'irq*, Arabic for "vein"), interspersed with gravel plains, salt flats, and scattered shrub communities[1,2]. Environmental conditions are characterized by large temperature fluctuations, irregular rainfall, and frequent dust storms[3]. Despite these harsh conditions, soil microbial communities also inhabit this extreme environment[4], where spatial variation in soil properties and temporal changes in environmental conditions may influence their composition and distribution. Understanding these communities requires observations that account for both the spatial heterogeneity of the landscape and its changing environmental conditions.

Previous studies of soils in Saudi Arabia have surveyed microbial communities along elevation gradients[5], responses to rainfall across soil compartments[6], and differences between bulk and rhizosphere soils across regions[7]. A study investigated microbial communities associated with desert plants[8]. Repeated sampling across sites, soil compartments, and time points provides an opportunity to extend this work by examining microbial communities across both spatial and temporal scales. Integrating data at such a scale introduces data management requirements as it becomes necessary

to preserve the relationship between sampling locations, collection events, laboratory analyses, and environmental measurements.

Several resources provide models for organizing environmental and microbiome observations. Multi-campaign initiatives, such as Tara Oceans, link campaigns, stations, and sampling events to environmental and sequence data[9]. MIxS/MIMARKS specify contextual metadata for environmental sequencing data[10]. The National Microbiome Data Collaborative promotes the use of persistent identifiers and shared schemas to connect microbiome datasets[11]. Knowledge graphs, including KG-Microbe[12] and MetagenomicKG[13], provide complementary approaches for integrating taxonomic, functional and environmental information. Building on these principles, this resource provides a linked representation of field samples, sampling events, laboratory processes, and environmental measurements from a repeated survey of the Rub' al-Khali.

The resource covers five campaigns between 2023 and 2025 along a 1,043 km transect across the Saudi Rub' al-Khali. Repeatedly sampled soil specimens connect amplicon profiles and shotgun metagenomes to a catalogue of 5,438 species-representative metagenome-assembled genomes. Field observations, archived-soil pH, X-ray fluorescence and climate data provide environmental context. Shared specimen identifiers connect these assays while retaining their different sampling coverage and processing histories. The Data Records section and the machine-readable cohort manifest distinguish the released biological products from the subsets used in each companion analysis.

To improve interoperability and enable integrated queries across heterogeneous data sources, the resource is represented using established ontologies and standardized identifiers. Sites, samples, processes, qualities, and measurement values are represented using the Semanticscience Integrated Ontology (SIO)[14]. This representation explicitly links sequencing results and environmental measurements to the samples and collection events they derive from, retaining their connections to the underlying data-generation processes. Chemical entities are encoded with ChEBI[15], environmental classifications with ENVO[16], and taxonomic identifiers with NCBI Taxonomy[17]. The resulting linked RDF modules support queries across sampling campaigns, samples, microbial communities, and environmental observations while preserving the context in which each observation was generated.

Technical validation assesses the consistency and integrity of the resource, including reconciliation with source data, measurement inclusion, identifier mappings, and structural constraints. Query-result comparisons assess whether the represented links recover the expected source observations. The reported scope distinguishes the ontology constructs and modules processed by the ELK reasoner[18]. The portal (https://rubalkhali.science/) provides site-level pages and a SPARQL endpoint for querying the resource. The Data Records section describes the available datasets, and Data Availability specifies their access conditions. By preserving specimen identity across amplicon profiles, shotgun metagenomes, refined genomes, gene calls, and geochemical measurements, this resource provides the shared foundation for companion manuscripts on landscape-scale and compartment community replacement

and encoded persistence traits[19] and on shallow-subsurface archaeal unknown and export-enriched protein repertoires[20].

2 Methods

Figure 1 summarizes the data products and their provenance; Table 1 collects the principal software versions. The sampling and assay protocols below specify how source observations enter those products.

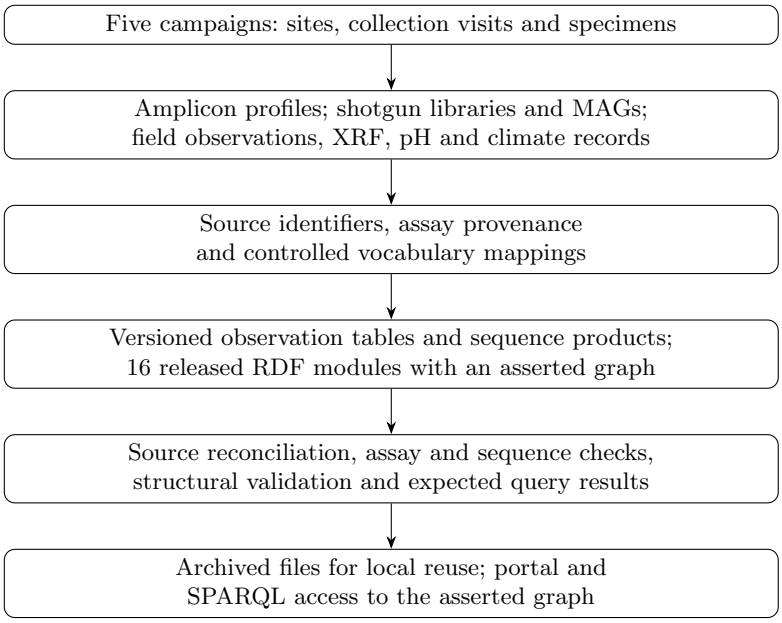


Fig. 1 Resource-generation and validation workflow. Associated sequence products and analytical tables retain their own provenance and accession requirements. The graph release contains the 16 modules enumerated in Data Records.

2.1 Sampling design and field observations

Five sampling campaigns covered a transect across the Rub' al-Khali starting from the western escarpment (19.15°N, 45.17°E) to the eastern Saudi Arabian border with Oman (20.71°N, 54.73°E). The 60 primary sampling sites span 1,015 km in a straight line and 1,043-km along the transect, calculated by summing great-circle distances between successive sites using an Earth radius of 6,371 km. Figure 2 shows the site coordinates and their altitudes.

In addition to the 60 primary sites, the survey included 10 named locations such as wells, a hot spring, and an oil-spill site. These locations were not part of the repeated sampling campaign. Comparisons across campaigns therefore focus on the 60 primary sites, and metadata links each observation to its sampling site, campaign, and sample.

Table 1 Principal software and reference versions. Version-specific execution settings and validation scope are described with the corresponding methods.

Stage	Software and references
Amplicon profiles	nf-core/ampliseq v2.14.0; DADA2; SILVA v138.2
Predicted functions	PICRUSt2 v2.4.1
MAG generation	MEGAHIT v1.2.9; SemiBin2 v2.2.1; DAS Tool v1.1.7; CheckM2 v1.1.0
MAG taxonomy	GTDB-Tk v2.7.2; GTDB R232
MAG annotation	Bakta v1.12.0 / database v6.0; Pyrodigal; eggNOG-mapper v2.1.13 / eggNOG v5.0.2; DIAMOND v2.0.15
MAG coverage	Bowtie2 v2.5.1; SAMtools v1.16.1; CoverM v0.6.1 and v0.7.0
Graph generation	OWL API v5.1.20; reference ontologies in Table 8
Graph validation	Jena ShEx v4.10.0; ELK v0.4.3 with OWL API v4.5.26

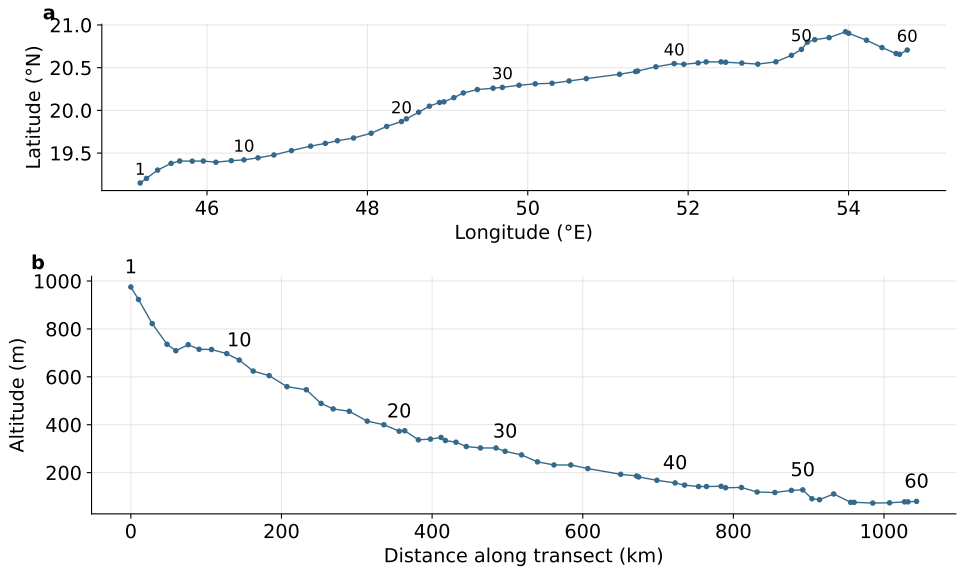


Fig. 2 Primary sampling transect across the Rub' al-Khali. (a) coordinates of sites 1–60. (b) altitude and cumulative great-circle distance.

The campaigns span different seasons: late winter 2023 (Trip 1), summer 2023 (Trip 2, truncated to 8 sites due to extreme temperatures), winter 2024 (Trip 3), summer 2024 (Trip 4), and autumn 2025 (Trip 5). Table 3 summarizes sites visited during each campaign.

At each site, we sampled three soil compartments: bulk surface soil, bulk subsurface soil (nominal 5–10 cm depth), and subsurface soil collected near roots (nominal

5–10 cm depth, hereafter "root-adjacent"). We collected three field replicates per compartment using sterile tubes and sterile plastic spoons, or three tubes taped together. Exact tube diameters, surface sampling depth, individual bulk subsurface depths and distance from roots were unrecorded. We stored samples at −20°C in the field in an Engel MR040F portable freezer, and transferred samples to an −80°C freezer in the laboratory at the end of each campaign. Table 2 summarizes soil samples collected by campaign and compartment.

Table 2 Soil samples collected and amplicon profiles obtained by campaign and compartment. Control samples are excluded.

	Surface		Deep		Root-adjacent		Total	
Trip	Collected	Profiles	Collected	Profiles	Collected	Profiles	Collected	Profiles
1	192	114	192	110	192	106	576	330
2	24	6	24	8	24	13	72	27
3	180	155	180	141	180	154	540	450
4	180	60	180	59	180	58	540	177
5	180	29	180	55	180	133	540	217
Total	756	364	756	373	756	464	2,268	1,201

2.1.1 Field observations

We recorded air temperature, atmospheric pressure, and relative humidity during site visits, although coverage varied between campaigns (Table 3). Temperature is reported in degrees Celsius, pressure in pascals, and relative humidity in percent. The field logs record pressure in hectopascals (millibars); we multiplied these values by 100 and assert them in pascals (`obo:UO_0000110`) in the graph. Trip 2 recorded only temperature, with eight measurements spanning 34.5–41.9°C. Field observations entered the dataset only when their site identity and measurement range could be established. Additionally, field researchers classified each site by an observed environmental feature (*e.g.* sand dune, saline pan, gravel), and we mapped these features to ENVO [16].

Table 3 Campaign overview and field measurement coverage.

Campaign	Season	Sites visited	Field measurements
Trip 1	Late Winter 2023	1–60; 4 named; 15 auxiliary*	Temp, Press, Hum.
Trip 2	Summer 2023	1–8	Temp.
Trip 3	Winter 2024	1–60; 5 named	Temp, Press, Hum.
Trip 4	Summer 2024	1–60	Temp, Hum.
Trip 5	Autumn 2025	1–60; 2 named	Temp, Press, Hum.

*The 15 observations dated March 2023 represent Trip 1 auxiliary/revisit observations.

2.2 DNA extraction and amplicon sequencing

Genomic DNA was extracted from 0.9–1.2 g of soil using either the QIAGEN PowerLyzer PowerSoil Kit (catalogue no. 12855-50) or the QIAGEN DNeasy PowerSoil Pro Kit (catalogue no. 47014). The extraction kit and method were recorded for samples when available. Frozen soil samples were thawed on ice and processed immediately. Before homogenization, samples were incubated at 65°C for 10 min with agitation at 1,200 rpm, then homogenized using a PowerLyzer 24 Homogenizer at 4,000 rpm for one 45-second cycle. The remaining purification steps were performed either manually or using a QIAcube Connect instrument, following the manufacturer's protocol. DNA was eluted in 50 µL of the elution buffer supplied with each kit, and stored at −20°C.

Extracted DNA was quantified using the Qubit dsDNA High Sensitivity assay, with fluorescence measured using a Varioskan LUX multimode plate reader (Thermo Fisher Scientific). Extract concentrations and processing information were recorded per sample.

The V3–V4 region of the bacterial 16S rRNA gene was amplified using the Bakt_341F (`CCTACGGGNGGCWGCAG`) and Bakt_785R (`GACTACHVGGGTATCTAATCC`) primers[21], both carrying Illumina overhang adapters. First-round PCR used TaKaRa Ex Taq DNA Polymerase (catalogue no. RR001B). Each 50-µL reaction contained 5 µL of 10× buffer with Mg^{2+}, 4 µL of dNTP mix (2.5 mM each), 2 µL of each primer (10 µM), 0.25 µL of polymerase (5 U/µL), 5 µL of DNA and 31.75 µL of nuclease-free water, giving final concentrations of 0.4 µM of each primer, 0.2 mM of each dNTP and 1.25 U of polymerase. Cycling conditions were 95°C for 3 min; 35 cycles of 95°C for 10 s, 65°C for 30 s and 72°C for 1 min; and a final extension at 75°C for 5 min.

PCR products were checked on a 1% agarose gel for an expected amplicon size of 500 bp. Samples without a visible product were re-amplified with 8 µL and then 12 µL of DNA, with DNA input set by volume because most extracts had low or undetectable concentrations. Samples that failed all three attempts were recorded as failed.

Dual indices and Illumina sequencing adapters were added in a second-round PCR using the QIAGEN Multiplex PCR Kit (catalogue no. 206143), following the Illumina 16S Metagenomic Sequencing Library Preparation protocol[22]: initial activation at 94°C for 15 min; 8 cycles of 94°C for 30 s, 55°C for 30 s and 72°C for 30 s; final extension at 72°C for 5 min.

Indexed libraries were purified using Agencourt AMPure XP beads (0.8:1 bead to sample volume), quantified with the Qubit dsDNA High Sensitivity assay, and checked for fragment size on an Agilent 4200 TapeStation (D5000 ScreenTape). Libraries were pooled in equimolar amounts, calculating molarity from concentration and fragment size. Pools were sequenced on an Illumina NovaSeq 6000 with an SP flow cell and a 500-cycle kit (2×250 bp paired-end reads). Library preparation was performed in the research laboratory and sequencing at the KAUST Bioscience Core Lab.

2.3 Laboratory controls and paired PMA treatments

Controls cover different processing stages, following reporting guidance for low-biomass samples[23] (Table 4). To control for contamination introduced during DNA extraction, extraction blanks were included with nuclease-free water as input instead of soil, and

processed in the same way as soil samples. Trips 1–3 used one extraction blank per extraction kit, and Trips 1–2 included shared extraction days; Trips 4–5 used one blank per trip-specific extraction-day batch. For Trip 4, we sequenced six extraction blanks (SEB, EB1–EB5); the blanks of three of the nine extraction batches were not sequenced. For Trip 5, we sequenced 18 extraction blanks (EB1–EB18), 17 of which have confirmed batch links to 217 biological profiles, whereas EB18 has no frozen extraction-day or co-extracted-sample mapping and is retained for characterisation only. The mapped blanks cover 95 canonical biological profiles for Trip 4 and 217 for Trip 5. These batch memberships are recorded separately from the campaign labels.

PCR negative controls comprised two types. No-template controls (NTCs) contained PCR reagents and nuclease-free water instead of DNA, and reagent-only controls per amplification plate that contained reagents but neither water nor DNA. Reagent-only controls had low or undetectable DNA concentrations and TapeStation profiles of only adapter dimers, so they were not sequenced. We sequenced six NTCs that were prepared with Trip 5 samples.

Positive standards assess recovery of expected genera. The Trips 1 and 2 protocol used the ZymoBIOMICS HMW DNA Standard (D6322), which contained purified DNA and was added at the PCR step. D6322 contains seven bacterial genomes (14% of genomic DNA each) and one yeast (2%). The Trip 3 protocol used the ZymoBIOMICS Microbial Community Standard (D6300) which contains whole cells and was therefore included from the extraction step. D6300 contains eight bacteria (12% of genomic DNA each) and two yeasts (2% each). These proportions are the manufacturer's values. Two of each type of positive standard were amplified and sequenced.

Table 4 Laboratory controls by campaign.

Trip	Extraction blanks	PCR negative controls	Positive standards
1	One per kit; PowerSoil Pro blank sequenced	Nuclease-free water	
2	One per kit	NTC-2	Two D6322 (DNA)
3	One per kit	PCRCtrl (index 728)	Two D6300 (whole cells)
4	One per batch; 6 of 9 sequenced	Not sequenced	No positive control
5	18 sequenced; 17 batch-linked	Six NTCs sequenced	D6300, shotgun sequenced

DNA extracted from soil includes DNA from intact cells as well as extracellular DNA and DNA from dead or damaged cells[24]. Propidium monoazide (PMA) is a dye that enters only cells with damaged membranes and binds to free DNA, and prevents amplification by PCR when exposed to light. The paired-treatment dataset records 16S rRNA gene profiles from PMA-treated and untreated aliquots of a subset of Trip 5 samples, with shared preparation and exposure conditions.

The PMA dataset includes three soil samples collected from campsites during Trip 5. The source material comprised root-adjacent soil from campsite C1, and

root-adjacent and surface soil from campsite C2. From each collection tube, three subsamples were taken and each suspended in phosphate-buffered saline (PBS, pH 7.4) at a soil-to-buffer ratio of 1:10 (w/v). The three samples from each collection tube are technical replicates of the treatment. Each sample was split into a treated aliquot that included 50 µM PMAxx, and an untreated aliquot, giving nine treated/untreated pairs.

Both aliquots followed the same exposure procedure: they were kept in the dark for 10 min to allow the dye to enter the cells, and then exposed to 450 nm light for 35 min at a distance of 15–20 cm from the light source, and turned every 15 min to ensure even exposure. Aliquots were stored at −20°C, and DNA was extracted from 860 µL of each aliquot with the PowerSoil Pro kit.

2.4 Amplicon profiles and predicted functions

Raw amplicon reads were processed with nf-core/ampliseq (v2.14.0)[25,26], using DADA2[27] for error correction and denoising, and SILVA (v138.2)[28] for taxonomic assignment. This resulted in 351,472 amplicon sequence variants (ASVs), and included 1,247 biological samples and 24 control samples. Quality filtering excludes the 24 control profiles and ten biological profiles with fewer than 1,000 reads, including one `T1Dr1` run with 934 reads. This leaves 1,237 profiles, including 1,227 profiles from sites 1–60, and 1,186 deduplicated specimen profiles for the primary transect; the complete table is distributed unfiltered. The container identity and full denoising parameters are included in the repository.

Gene family and pathway abundances were predicted from the 16S rRNA data with PICRUSt2 v2.4.1[29]. Two executions used the Trips 1–4 ampliseq output and the Trip 5 output (330,830 ASVs). KEGG Orthology (KO) prediction used the phylogenetic independent contrasts method, edge exponent 0.5, and random seed 100. Enzyme Commission (EC) and KO abundances required 100 reads, and pathway abundances were inferred from the EC predictions.

2.5 Soil chemistry and climate observations

Soil elemental composition was measured by X-ray fluorescence (XRF) using an Olympus Vanta VMR-XXX-G3-E handheld analyser (serial no. 841307), both in the field during Trip 5 and in the laboratory on archived soil samples from all campaigns. The field log contains 106 entries across 59 sites. For field measurements, the instrument was used in Geochem (3-Beam) mode with NORMAL settings, and 71 complete measurement sessions from 58 sites were retained, including repeated sessions at nine sites. In the laboratory, we analyzed 725 samples with the instrument in Fast Mode with the Fast Screening-He8mm setting. Instrument-reported elemental and oxide values are stored as separate variables with undocumented units.

To assess agreement between field and laboratory measurements, the comparison paired the median of repeated field sessions at each site with the corresponding Trip 5 laboratory measurements at the 58 sites measured using both methods. This comparison assesses agreement in site-level rankings, with correlations reported in Technical Validation.

The mapping cross-referenced the 93 instrument reporting channels against ChEBI r247 (3 December 2025) and PubChem properties (retrieved 23 July 2026) [15,30]. The mapping identified corresponding ChEBI entries for 79 channels and PubChem entries for 91 channels. Verification covered identifier validity, names, chemical formulae, charge, and whether each variable represented an element or an oxide.

Archived-soil pH measurements cover all five campaigns. Dried soil was sieved to below 2 mm; 1 g was mixed with 2.5 mL of 0.01 M $CaCl_2$. Each suspension was mixed for 30 min in a HulaMixer and left to stand for another 30 min at room temperature before measuring pH.

The pH measurements used an accumet AB315 meter with a Fisher Scientific electrode (cat. no. 13-620-112) calibrated at the start of each session, with the electrode slope and the reading of the pH 10 buffer recorded for every session. The electrode was rinsed with Milli-Q water between measurements. Of 767 recorded measurements, 36 were quarantined because of date provenance and 19 because of quality control. The retained dataset contains 712 measurements with a recorded electrode slope between 95% and 102% and a pH 10 buffer reading within 0.1 units of its nominal value (Table 5). All 156 admitted Trip 4 pH measurements used archived field replicate 2. Each admitted measurement is linked to its physical specimen and retains its original assay identifier. The ecology analysis uses the same measurement values.

Table 5 Soil pH measurement coverage.

Campaign	Surface	Subsurface	Root-adjacent	Total
1	52	56	49	157
2	5	6	12	23
3	54	50	52	156
4	50	52	54	156
5	30	57	133	220
Total	191	221	300	712

Historical climate data were retrieved for each sampling site from the Open-Meteo archive API [31] using geographic coordinates, yielding both monthly and daily climate observations. The monthly dataset contains 12,936 values that cover temperature, total precipitation, rainfall, and mean relative humidity from January 2022 to January 2026. The daily dataset covers 1 January 2022 to 1 February 2026 and covers temperature, total precipitation, and rainfall. The monthly and daily products come from two separately configured acquisitions; the monthly acquisition ends on 20 January 2026, so January 2026 is a partial month in the monthly data.

Daily precipitation data were also retrieved from NASA POWER PRECTOTCORR for sites 1–60 from 1 January 2022 to 31 December 2025, which yielded 87,660 records.

2.6 Shotgun metagenomics and genome catalogue

Shotgun metagenomic sequencing complements the amplicon profiles with genome assemblies and gene annotations from the same sampling campaigns. The read-bearing library inventory contains 551 libraries across 61 site codes (59 primary transect sites, comprising all primary sites except site 56, plus two Trip 1 named locations at sites 62 and 64): 210 from Trip 1, 37 from Trip 2, 128 from Trip 3, 85 from Trip 4 and 91 from Trip 5. Library metadata records the campaign, source site, soil compartment and replicate. The libraries were sequenced as 554 paired-end runs, 221 on an Illumina NovaSeq 6000 and 333 on an Illumina NovaSeq X; three libraries were sequenced on both instruments. Libraries for 81 runs were prepared with the NEBNext Ultra II DNA Library Prep Kit and for 473 runs with the QIAseq FX DNA Library UDI Kit. The shotgun study is registered at ENA under `PRJEB125153` (secondary accession `ERP204894`).

Genome reconstruction combined MEGAHIT v1.2.9 assembly [32] with single-sample and self-supervised SemiBin2 v2.2.1 binning [33], followed by DAS Tool v1.1.7 refinement [34]. MEGAHIT produced contigs of at least 500 bp for 549 of the 551 libraries; libraries 51Sr1 and F5Dr1 yielded no contigs. CheckM2 v1.1.0 [35] estimated completeness and contamination for the refined bins, and 12,434 refined bins reached at least 50% completeness with at most 10% contamination. The bins were dereplicated with dRep [36], using CheckM2 quality estimates, a minimum completeness of 50%, a maximum contamination of 10%, and a minimum genome length of 50,000 bp. The recorded clustering parameters specify 90% primary ANI and 95% secondary ANI, with fastANI for secondary comparisons and a minimum alignment coverage of 0.1 using dRep's `larger` coverage mode. The resulting catalogue contains 5,438 representative identifiers. The provenance table links each identifier to its source library, binning strategy and bin, and records quality estimates from the refined-bin CheckM2 report.

Taxonomic assignments for the refined representatives used GTDB-Tk v2.7.2 with GTDB release R232 [37]. The two native classification tables contain 4,847 bacterial and 591 archaeal representatives, and their genome identifiers match the representative catalogue exactly. A frozen historical analytical taxonomy snapshot (`EQ-TAX-LEGACY-ANALYTICAL-v0`) used by earlier companion genus/order turnover and proteome scripts shares the same 5,438 genome identifiers and identical CheckM2 quality metrics, as well as the exact same 237 Thermoplasmatota genome identifiers, but differs from the canonical GTDB R232 strings on 1,898 of the 5,438 genomes (zero at domain, 247 at phylum, 34 at class, 220 at order, 511 at family, 1,124 at genus and 46 at species; 42 differing strings among the 237 Thermoplasmatota genomes, where canonical GTDB R232 assigns four named orders plus two unassigned and 19 named genera plus 135 unassigned versus four named orders plus zero unassigned and 14 named genera plus 139 unassigned in the frozen snapshot). The two taxonomy snapshots are linked 1:1 in the packaged reconciliation table and never conflated. The annotation workflow detected rRNA genes with barrnap v0.9 and tRNA genes with tRNAscan-SE v2.0.12 [38], and assigned each representative a quality tier following the MIMAG standard [39]. Of the 5,438 representatives, 126 meet the high-quality criteria, which require more than 90% completeness, less than 5% contamination, the 23S, 16S and 5S rRNA genes and tRNAs for at least 18 of the 20 amino acids. Another

5,311 meet the strict medium-quality criteria of at least 50% completeness and less than 10% contamination, giving 5,437 strict HQ+MQ representatives deposited as ENA MAGs. The remaining representative, `48Dr2__self_SemiBin_25` (completeness 64.61%, contamination 10.0%), sits at the 10% contamination boundary; it is classified as below-medium under strict MIMAG criteria and deposited in ENA as a binned metagenome, while operational dRep analyses using the inclusive $\leq 10\%$ contamination threshold include all 5,438 species representatives (126 high-quality and 5,312 non-high-quality representatives).

Two distinct gene-calling products are maintained across the 5,438 representatives (Table 9). In the canonical release annotation product (snapshot identifier `EQ-PROT-BAKTA-v1.0.0`), the genes of each representative were annotated with Bakta v1.12.0[40] and its full database v6.0, predicting coding sequences and translated proteins alongside rRNA and tRNA genes. eggNOG-mapper v2.1.13[41] searched these predicted proteins against the eggNOG v5.0.2 database with DIAMOND v2.0.15 to assign functional annotations, KEGG orthologs (KOs), Pfam domains and CAZy families (counts in Table 9). Per-genome summaries count the predicted proteins, the proteins with an eggNOG assignment and the proteins with at least one KO, and list the genome–feature pairs that occur in the annotations. In the companion analysis protein product (snapshot identifier `EQ-PROT-PYRODIGAL-REP-v1.0.0`), Pyrodigal predicted 20,451,986 proteins associated with the same 5,438 representative genome identifiers (reconciled against 20,379,930 bin-mapped assembly protein calls drawn from 753,305,533 whole-assembly protein predictions), supplying the sequence universe for downstream functional and archaeal proteome analyses[19,20]. Direct locus-by-locus equivalence between the 19,186,792 Bakta predictions and the 20,451,986 Pyrodigal predictions has not been established; matching genome identifiers do not establish identical underlying FASTA sequence contents or retained contig coordinates, and the two protein sets are maintained as separate products. Companion genome-resolved KO projections likewise use a frozen analytical KO matrix (`EQ-KO-LEGACY-ANALYTICAL-v0`, 11,386 distinct KOs across the 5,438 representatives and 8,860 KOs shared with PICRUSt2) that is recorded separately from `EQ-PROT-BAKTA-v1.0.0`.

Four read-mapping products with distinct snapshot identifiers record representative abundance and detection across shotgun libraries. First, to estimate per-library coverage in the canonical depth/breadth product (`EQ-COVERM-REFINED-v1.0.0`), we mapped the reads of all 551 libraries to a single reference that contains every contig of the 5,438 representatives. Each contig identifier was prefixed with its genome identifier, so that contigs with the same assembler-assigned name in different genomes remain distinct. Bowtie2 v2.5.1[42] aligned the first 200 million read pairs of each library in `--very-sensitive` mode and retained aligned reads only; SAMtools v1.16.1[43] sorted the alignments. CoverM v0.6.1[44] in genome mode counted reads aligned with at least 95% identity and reported, for every genome, the mean depth, the covered fraction, the trimmed mean depth, the number of mapped reads and the genome length. Second, a CoverM v0.7.0 genome-mode relative-abundance matrix (`EQ-COVERM-REL-ABUND-v1.0.0`, retaining the unmapped read fraction across 573 matrix columns including 551 study libraries) defines the current genome-resolved

ecological cohorts[19]. Exact specimen joins and post-QC read-depth availability define these nested subsets (Table 6), while per-library admission flags and excluded sample identifiers are recorded in the packaged cohort accounting table. Third, an archaeal competitive-mapping product (`EQ-COVERM-ARCHAEA-COMP-v1.0.0`) maps 193 paired metagenomic profiles across 68 site-campaign visits and 36 sites against the 5,438 catalogue representatives combined with 301 external Thermoplasmatota reference genomes across 0%, 10%, 20% and 50% genome-breadth gates, with a `MAPQ` $\geq$ 20 filter evaluated as a separate sensitivity analysis[20]. Fourth, a historical Trips 1–3 pilot dataset (`EQ-COVERM-LEGACY-991-v0`) retains 150 CoverM profiles (58, 15 and 77 libraries from Trips 1–3) mapped against an earlier 991-genome draft catalogue, eggNOG-mapper v2.1.12[41] annotations yielding a 990-genome KO copy-number matrix, and six source tables for an initial MAG-derived versus PICRUSt2 comparison; this pilot product is preserved for audit reproducibility and superseded for full-transect analyses by the 5,438-representative products.

Table 6 Nested shotgun metagenomic cohorts derived from the 551 study libraries. Product key: A, canonical Bowtie2/CoverM v0.6.1 depth and breadth profile archive (snapshot identifier `EQ-COVERM-REFINED-v1.0.0`); B, CoverM v0.7.0 genome relative-abundance cohort matrix (snapshot identifier `EQ-COVERM-REL-ABUND-v1.0.0`).

Cohort	Count	Defining exclusion or intersection	Product
All study libraries	551	Read-bearing libraries across all five campaigns	A
Non-zero mapped	550	Excludes library with zero mapped reads (`F5Dr1`)	B
Core-transect	546	Excludes four Trip 1 named-site libraries (`62PRr2`, `64Dr3`, `64PRr2`, `64Sr2`); spans 59 primary sites	B
Core post-QC depth	545	Excludes core library lacking post-QC read-depth metadata (`F45PRr3`)	B
Core matched 16S	442	Exact specimen join to the 1,186-sample deduplicated 16S cohort across 58 primary sites (excludes sites 56 and 59; 8,860 shared KOs with PICRUSt2)	B
Matched post-QC depth	441	Intersects 442 matched cohort with complete post-QC read-depth metadata (excludes `F45PRr3`)	B

2.7 Knowledge representation

The resource is represented as a modular RDF knowledge base that connects sampling sites, physical samples, laboratory measurements, experimental procedures, sequencing datasets, and taxonomic assignments. Explicit relations connect observations to

their sources. This allows users to trace an observation to the sample, site, or process it describes, and query related data across different stages of the study (Figure 3).

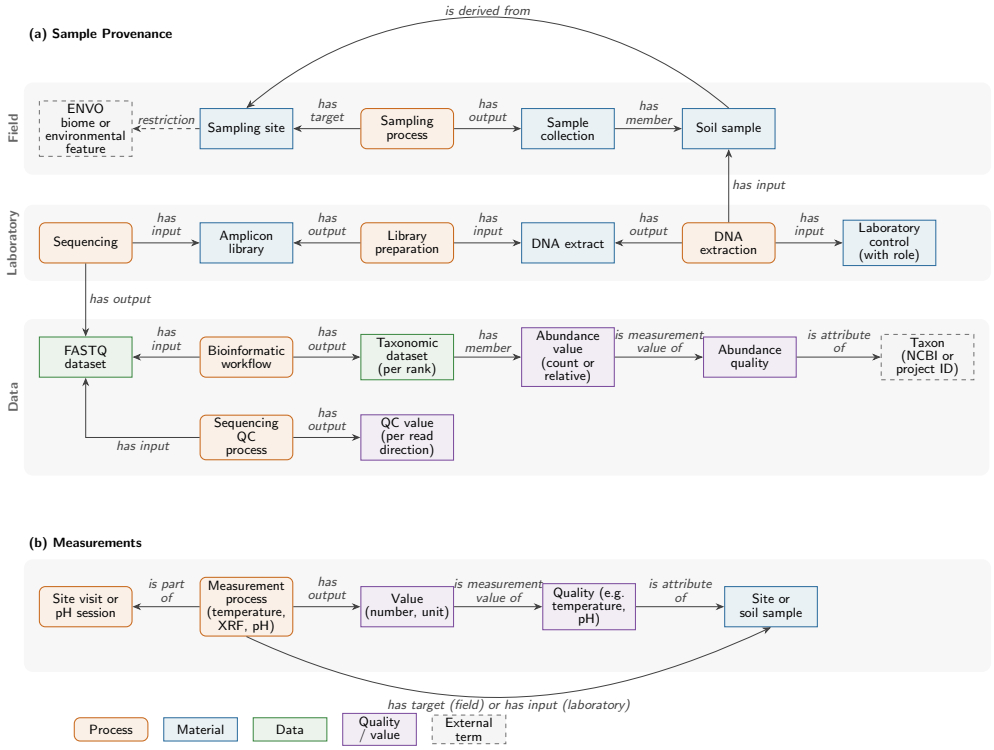


Fig. 3 Main patterns in the knowledge graph. (a) A sampling process targets a site and produces a collection of soil samples, each also linked directly to its site. Laboratory processes are connected through their inputs and outputs, from DNA extraction to sequencing; laboratory controls enter at the processing step whose role they test. A bioinformatic workflow converts each FASTQ dataset into taxonomic datasets at seven ranks, whose values quantify the abundance of taxa identified by NCBI Taxonomy or project identifiers. Sites are linked to ENVO classes through existential restrictions (dashed). (b) Each measurement process produces a value that quantifies a quality of a site or soil sample, and is part of a site visit or pH measurement session. Arrows point from subject to object of each relation. Protocols linked to each process are omitted for clarity.

The graph separates a shared terminology module, which defines the classes and relationships used to describe the data, from modules containing the actual study records. In ontology terminology, these correspond to the TBox and ABox, respectively. Shared identifiers connect the modules (Table 8), while the Semanticscience Integrated Ontology (SIO) provides general classes and relationships for representing entities, processes, and measurements [14]. The terminology module declares 333 classes (297 defined for this project and 36 from other ontologies), 20 object properties, and 35 datatype properties written in OWL [45].

To support interpretation and reuse, KG v3.0.1 provides English definitions for all 342 project terms, recorded using SKOS and IAO definition annotations. Claude Sonnet 5.5 drafted the definitions from term labels, axioms and dataset context, and OpenAI Codex assisted with automated review and editing. The release distributes the definitions, model provenance and review records. The terminology module imports four reference ontologies, and the observation modules are linked as parts of the dataset.

2.7.1 Representing provenance

The data model rests on two distinctions. First, physical material (a soil sample, a DNA extract), the process applied to it (extraction, sequencing), and the data those processes produce (a FASTQ dataset) are represented as separate entities connected by inputs and outputs. Second, a measured value is kept separate from the property it quantifies and from the entity that has the property. Together, these distinctions record where every value came from.

Established ontology terms describe shared entity types (Table 7). Because these identifiers are shared with other resources, the data can be queried and compared alongside other datasets. Original labels and values are retained with the mapped identifiers.

Table 7 External ontologies and reference resources used in the knowledge graph.

Resource	Used to describe
SIO[14]	Processes, inputs and outputs, values, qualities and the relations between them
PATO	Measured qualities
UO[46]	Units of measurement
ENVO[16]	Biomes and environmental features of sites
NCBI Taxonomy[17]	Taxa
ChEBI[15], PubChem[30]	Elements and oxides reported by XRF

Measurements are represented separately from the biological or environmental qualities they describe. For example, the temperature of a site is a property (quality) of that site (Listing 1). A measurement process carried out during a particular site visit produces a value for that property. The graph links the value to the quality, the quality to the site, and the measurement process to the visit. Because SIO provides relations in both directions between a value and its quality, queries can start from

either and retrieve each measurement together with its sampling context. This separates what was measured from how, where, and when it was measured, and allows the same quality to be measured repeatedly across visits without ambiguity.

Each value holds its numerical reading and, where documented, its unit. The ShEx shape assigned to each type of value specifies its datatype and unit requirements, which differ for counts, measurements with units and measurements whose units were not documented.

Listing 1 shows this representation for a field temperature measurement. The measurement process is part of a site visit, the visit targets the site, and the process produces a value of 20.7 °C, which quantifies the temperature of that site. Because SIO provides relations in both directions between a value and its quality, users can start a query from either and retrieve each measurement together with its sampling context.

Listing 1 Field temperature and its visit provenance.

```
@prefix rak: <https://rubalkhali.science/kb/RAK_> .
@prefix sio: <http://semanticscience.org/resource/SIO_> .
@prefix pato: <http://purl.obolibrary.org/obo/PATO_> .
@prefix uo: <http://purl.obolibrary.org/obo/UO_> .
@prefix xsd: <http://www.w3.org/2001/XMLSchema#> .

# Measurement process
rak:P000001 a rak:0000006 ; # type: field measurement process
    sio:000068 rak:3000001 ; # is part of -> site visit
    sio:000229 rak:4000001 . # has output -> measured value

# Site visit
rak:3000001 a rak:0000003 ; # type: site visit
    sio:000291 rak:1000001 . # has target -> sampling site

# Measured value
rak:4000001 a rak:0000010 ; # type: measurement value
    sio:000221 uo:0000027 ; # has unit -> degree Celsius
    rak:2000003 "20.7"^^xsd:double ; # numerical value
    sio:000215 rak:5000001 . # is measurement value of -> quality

# Quality of the site
rak:5000001 a pato:0000146 ; # type: temperature (PATO)
    sio:000011 rak:1000001 ; # is attribute of -> sampling site
    sio:000216 rak:4000001 . # has measurement value -> value
```

A taxonomic label can denote different taxa across databases and releases, and a taxon can be referred to by several labels[47]. To ensure consistent identification, the mapping retains SILVA assignments for each ASV and links each ASV to either a validated NCBI Taxonomy identifier or a project-specific identifier when no suitable NCBI Taxonomy identifier was found.

Each SILVA assignment describes the organism across seven ranks from domain to species, and may include an additional species prediction and a confidence score. The representation uses the species name from the rank assignment in the knowledge graph and retains the additional species assignment as an annotation. Normalization removes taxon-name prefixes and explicitly records missing and unclassified ranks.

For each SILVA assignment, the mapping procedure checks for a candidate NCBI identifier from a fixed NCBI Taxonomy release (OBO Foundry NCBITaxon r2025-03-13 built from NCBI Taxonomy 1 March 2025). A candidate was only accepted if the identifier exists in that release, its rank matches, its name matches the SILVA name or a documented synonym, and its parent taxa in NCBI match identifiers already accepted for the higher ranks of the same assignment. When no candidate met these criteria, the name was given a project-specific identifier, which refers to a unique combination of rank, name, and parent taxa. Users can query each ASV by its original SILVA assignment or by its NCBI or project identifier.

2.7.2 Modelling patterns and their implementation

The following description-logic expressions summarize the intended entity relationships[48]. The published project TBox declares class hierarchies and property axioms, including four property chains, but contains no explicit restriction nodes. The expressions below therefore describe modelling patterns; the released ShEx shapes specify which corresponding links are required in stored data. Each shape has the scope and limitations documented in Section 2.7.4. In the notation below, $\sqsubseteq$ means "is a subclass of", $\sqcap$ means "and", $\exists R.C$ means "has at least one relation R to a member of class C", and $\forall R.C$ means "all relations R point to members of class C". Section 5.2 gives the corresponding RDF examples.

A sampling process targets a site and produces a collection of samples, and each sample is also linked directly to its site.

$$\text{SamplingProcess} \sqsubseteq \text{Process} \tag{1}$$
$$\text{Sample} \sqsubseteq \text{MaterialEntity} \tag{2}$$
$$\text{SampleCollection} \sqsubseteq \text{Collection} \sqcap \forall \text{hasMember.Sample} \tag{3}$$
$$\text{SamplingProcess} \sqsubseteq \exists \text{hasTarget.Site} \sqcap \exists \text{hasOutput.SampleCollection} \tag{4}$$
$$\text{Sample} \sqsubseteq \exists \text{isDerivedFrom.Site} \tag{5}$$

Each field or laboratory step is recorded as the execution of a protocol, with the material or data it used and produced.

$$\text{ProtocolExecution} \sqsubseteq \text{Process} \tag{6}$$
$$\text{Protocol} \sqsubseteq \text{InformationContentEntity} \tag{7}$$
$$\text{ProtocolExecution} \sqsubseteq \exists \text{isSpecifiedBy.Protocol} \tag{8}$$
$$\text{ProtocolExecution} \sqsubseteq \exists \text{hasInput.Entity} \sqcap \exists \text{hasOutput.Entity} \tag{9}$$

A value quantifies the abundance of a taxon and belongs to a dataset produced by a bioinformatic workflow from a sequence dataset.

$$\text{BioinformaticWorkflow} \sqsubseteq \text{Process} \tag{10}$$
$$\text{TaxonomicDataset} \sqsubseteq \text{Collection} \tag{11}$$
$$\text{BioinformaticWorkflow} \sqsubseteq \exists \text{hasInput.SequenceDataset} \sqcap \exists \text{hasOutput.TaxonomicDataset} \tag{12}$$

$$\text{TaxonomicDataset} \sqsubseteq \forall \text{hasMember.AbundanceObservation} \tag{13}$$
$$\text{AbundanceObservation} \sqsubseteq \text{Quantity} \sqcap \exists \text{isMeasurementValueOf.AbundanceQuality} \tag{14}$$
$$\text{AbundanceQuality} \sqsubseteq \exists \text{isAttributeOf.Taxon} \tag{15}$$
$$\text{AbundanceQuality} \sqsubseteq \exists \text{hasMeasurementValue.AbundanceObservation} \tag{16}$$

Laboratory XRF measurements take a soil sample as input, whereas field measurements target a sampling site, and both produce concentration values.

$$\text{XRFProcess} \sqsubseteq \text{MeasurementProcess} \sqcap \exists \text{hasOutput.Concentration} \tag{17}$$
$$\text{LabXRFProcess} \sqsubseteq \text{XRFProcess} \sqcap \exists \text{hasInput.SoilSample} \tag{18}$$
$$\text{FieldXRFProcess} \sqsubseteq \text{XRFProcess} \sqcap \exists \text{hasTarget.SamplingSite} \tag{19}$$
$$\text{Concentration} \sqsubseteq \text{Quantity} \sqcap \exists \text{isMeasurementValueOf.AnalyteQuality} \tag{20}$$

DNA extraction, library preparation and sequencing are linked through their shared inputs and outputs, so that each sequence dataset can be traced back to its soil sample.

$$\text{DNAExtractionProcess} \sqsubseteq \text{ProtocolExecution} \tag{21}$$
$$\text{DNAExtractionProcess} \sqsubseteq \exists \text{hasInput.SoilSample} \sqcap \exists \text{hasOutput.DNAExtract} \tag{22}$$
$$\text{LibraryPreparationProcess} \sqsubseteq \text{ProtocolExecution} \tag{23}$$
$$\text{LibraryPreparationProcess} \sqsubseteq \exists \text{hasInput.DNAExtract} \sqcap \exists \text{hasOutput.AmpliconLibrary} \tag{24}$$
$$\text{SequencingProcess} \sqsubseteq \text{ProtocolExecution} \tag{25}$$
$$\text{SequencingProcess} \sqsubseteq \exists \text{hasInput.AmpliconLibrary} \sqcap \exists \text{hasOutput.FASTQDataset} \tag{26}$$
$$\text{FASTQDataset} \sqsubseteq \text{Dataset} \tag{27}$$
$$\text{FASTQDataset} \sqsubseteq \forall \text{hasMember.SequenceRead} \tag{28}$$

Quality metrics are attached to the FASTQ dataset they describe.

$$\text{SequencingQCProcess} \sqsubseteq \text{Process} \tag{29}$$
$$\text{SequencingQCProcess} \sqsubseteq \exists \text{hasInput.FASTQDataset} \sqcap \exists \text{hasOutput.QCMeasurementValue} \tag{30}$$
$$\text{QCMeasurementValue} \sqsubseteq \text{Quantity} \sqcap \exists \text{isMeasurementValueOf.QCQuality} \tag{31}$$
$$\text{QCQuality} \sqsubseteq \exists \text{isAttributeOf.FASTQDataset} \tag{32}$$

2.7.3 OWL and RDF implementation

The modules are generated with Groovy scripts using OWL API v5.1.20[49]. Keeping the shared terminology separate from the study data allows each to be updated and checked independently[50]. The asserted graph is stored in OpenLink Virtuoso[51] which supports SPARQL queries over the whole resource. Release KG v3.0.1 serves the asserted triples.

Some facts follow from the asserted statements without being stated explicitly; for example, a sample of a subclass of soil sample is also a soil sample. A separate

workflow derives such facts with a rule program v1.5.0 of positive RDF rules and stores them in a separate graph, so that recorded statements can be distinguished from those that are derived. The rules cover equivalent classes and properties, subclass and subproperty hierarchies, instance types, property domains and resource-valued ranges, and inverse, symmetric, and transitive properties. For property chains, one rule is generated for each non-empty finite chain in the loaded schema, including chains with inverse properties. Further rules handle value restrictions, existential restrictions (classified from existing links and typed targets), universal restrictions (applied to existing targets), intersections, union membership and enumerated classes. The rules do not create new RDF terms. The schema audit lists the imported ontologies and the OWL constructs not covered by the rules.

The implementation computes the least fixed point of these rules, with a complete pass adding no triples as the stopping criterion. Its run metadata and clash diagnostics define the scope of the inferred output. Usage Notes describes execution and completion checks separately from the asserted release.

2.7.4 Validation

The generated graph was checked for structural integrity with Shape Expressions (ShEx) and for logical consistency with an OWL reasoner. These checks assess required graph structure and the logical consequences within the supported reasoning scope. Results are reported in Section 4.

ShEx[52] is a schema language that checks whether selected entities have the required types and properties, how many values each property has and whether the values have the correct datatype. The resource defines shapes for all main entity types: sampling sites, soil samples, DNA extracts, sequencing datasets, measurements, quality-control records, laboratory control materials, control roles, laboratory batches, sequences detected in controls, soil pH measurements, values and sessions, and taxonomy mappings. For example, the sample shape requires a replicate type, a label, a source identifier and a link to the site from which the sample was derived, and the DNA concentration shape checks the numerical value, the unit (ng/µL) and the links to the measured quality and process.

The shapes have known limits. The site shape checks that a geometry is present, but not that it is valid WKT with correct coordinate syntax. The XRF-process shape requires an input or target, but checking that it belongs to the correct class requires an additional check. The shapes are open, so entities may have properties beyond those declared, and additional types are accepted regardless of how they were asserted. Test data that deliberately violate the shapes confirm that the failures reported in Technical Validation are detected.

A Shape Map assigns each set of entities (e.g. all sampling sites) to its shape. During deployment, each newly generated data module is checked, and deployment stops if a required property is missing (e.g. the coordinates of a site) or a value has the wrong datatype. Validation used the Apache Jena ShEx library v4.10.0, called from Groovy scripts. The multi-gigabyte taxonomy data are too large for this in-memory check; instead, the file is parsed in full with an independent Turtle parser and checked with

a streaming validator that covers every link between classification process, dataset, quality and value.

Logical consistency was checked with the ELK reasoner v0.4.3 [18] through OWL API v4.5.26, after merging the modules that ELK can process into a single ontology. ELK supports only part of the OWL language, and the modules also contain universal restrictions, inverse properties and other constructs that ELK does not handle. The result therefore applies only to the supported subset; an OWL profile report and a checked projection of the ontology would be needed to define its logical scope precisely. The taxonomy data and large reference modules were excluded.

2.8 Genome-resolved trait annotation

We profiled all 5,438 representative genomes with microTrait [53], using their checksum-verified nucleotide sequences. This product, `EQ-MICROTRAIT-REP-v1.0.0`, uses fresh Prodigal v2.6.3 calls in metagenomic mode and retains its own gene identifiers and provenance, separately from the Bakta and companion Pyrodigal protein sets. Profile searches used the pinned microTrait models and the selected dbCAN V14 glycoside-hydrolase models, with the published profile thresholds and logical trait rules. The package records the exact microTrait and gRodon source revisions, database assets, software environment and execution parameters.

The outputs comprise model-hit and rule matrices, trait matrices at three levels of aggregation, and per-genome growth and optimum-temperature predictions with applicability flags. Trait values count satisfied rules within a category, with binary categories capped at one. Predicted growth and temperature values describe model-derived genomic properties; their interpretation retains genome completeness, fragmentation and archaeal applicability information. Of 1,831 model identifiers required by the rules, 1,829 were assayed. The release lacked the `aphA` and `phoN` models. For the phosphatase rule combining these models with `phoA` and `phoD`, the interpretation-safe table records a positive call when either assayed model supports it and otherwise records an unassessed value. The unmodified software output is retained alongside this coverage-aware table.

2.9 Reproducible workflow

A Nextflow [54] workflow generates the knowledge graph modules, runs the validation checks, and computes the statistics reported in this manuscript. Configuration profiles are available for running the workflow locally or on a computing cluster. Sample processing (Section 2) was run separately with its own parameters and software environments.

The workflow produces each versioned release of the knowledge graph by generating the data modules, comparing their checksums with those listed in the release manifest, loading them into a new database, and checking the results of test queries against the source data. The public service is switched to the new release only after these checks pass. The release manifests list the fixed input files and the generated modules along with their checksums.

3 Data records

KG v3.0.1 contains 45,707,660 asserted triples from 16 modules and is archived at Zenodo[55]. The resource contains the modules listed in Table 8, the combined asserted graph, and a reproducibility package with code, term definitions and validation evidence. The modules use RDF/XML and Turtle, and the combined graph uses compressed N-Quads. Shared site, sample and process identifiers connect observations across modules. The deposit also includes an inventory with checksums.

Table 8 The 16 RDF modules included in KG v3.0.1.

Module	Content
Project terminology	Base ontology (TBox)
Sampling sites	Sampling sites with coordinates and biome annotations
Samples	Soil and plant sample individuals with compartment types
Environmental observations	Field environmental measurements and climate records
XRF measurements	XRF measurement values from the field and laboratory workflows
DNA extracts	DNA extracts with concentration data
Sequencing runs	Sequencing run records linked to ENA accessions
Sequence quality	Sequencing quality metrics (GC content, duplicate rate, sequence counts)
Laboratory controls	Laboratory-control materials, roles, batches and sequence occurrences
Soil pH	Archived-soil pH measurements, sessions and specimen links
Taxonomic observations	Taxon observations over feature profiles
Taxonomic vocabulary	Vocabulary supporting canonical taxonomic mappings
SIO	Semanticscience Integrated Ontology v1.59
ENVO	Environment Ontology, 20 October 2025
PATO	Phenotype and Trait Ontology, 14 May 2025
Units Ontology	Units Ontology, 25 May 2023

Table 9 distinguishes the biological products from their graph metadata. The sample catalogue links collection visits, specimens, extractions, libraries and runs; these entities retain separate identifiers because a specimen can yield several assays or runs. The data dictionary specifies types, units and missing-value conventions; the climate inventory records acquisition windows and variables; and the file inventory records sizes, checksums, licensing and retrieval information. Archived-soil pH measurements and their Trip 4 replicate-2 specimen reconciliation are distributed under version identifier `EQ-PH-SHARED-v1.0.1`.

A machine-readable cohort and product manifest assigns manuscript and reproducibility snapshot identifiers to distinguish the canonical released resource from companion analytical products and historical pilots until formal repository deposition. These snapshot identifiers link the canonical 5,438-representative genome catalogue, the separate Bakta (`EQ-PROT-BAKTA-v1.0.0`) and Pyrodigal (`EQ-PROT-PYRODIGAL-REP-v1.0.0`) gene-calling products, and the four CoverM read-mapping products (`EQ-COVERM-REFINED-v1.0.0`, `EQ-COVERM-REL-ABUND-v1.0.0`, `EQ-COVERM-ARCHAEA-COMP-v1.0.0` and `EQ-COVERM-LEGACY-991-v0`) to the exact specimen and library cohorts reused by the companion studies[19,20]. The paired PMA profiles and the historical Trips 1–3 991-genome pilot analysis (150 CoverM profiles, a 990-genome KO matrix and six source tables for the initial MAG-derived versus PICRUSt2 comparison) retain their own reference definitions and analysis provenance. The microTrait trait profile (`EQ-MICROTRAIT-REP-v1.0.0`) provides genome-resolved trait matrices and continuous predictions as a separate annotation product, with compact tables and provenance in the reproducibility package and larger intermediates retained on institutional storage. This annotation product is separate from the archived KG v3.0.1; its bulk deposit is included in the pending archive described in Data Availability.

Tabular data use delimited text formats while larger collections are compressed.

3.1 FAIR identifiers and reuse metadata

The version-specific Zenodo DOI, ENA accessions and project-domain IRIs support discovery and identification of the resource[56]. The archived modules provide fixed inputs for local access, while the public endpoint provides query access to the deployed asserted graph. The release manifest identifies files by checksum and the data dictionary defines fields, units and missing-value conventions. Data Availability records the access conditions for sequences and associated products.

Interoperability is supported by RDF and shared SIO, ENVO, ChEBI, PATO, UO and validated taxonomic identifiers. Original labels and classifier lineages remain linked to their mapped identifiers, preserving the evidence used to interpret mappings. Specimen, process and measurement provenance, term definitions, validation evidence and scoped licences support reuse. Imported vocabularies retain their own licences; project-authored metadata and software have the licences stated in Data Availability.

4 Technical validation

Validation assesses source-to-release reconciliation, assay controls, sequence-file integrity and the accuracy of represented links. Structural constraints and logical checks assess the graph within their stated coverage. Expected query results provide an independent comparison with source observations. For every check, the validation records identify input files by checksum and the software versions used (Section 3).

4.1 Source reconciliation and measurement validation

The archived evidence workflow produced a specimen catalogue and reconciliation summaries for specimens, extracts, libraries, runs and community profiles. The 2,516 eligible source rows reconcile exactly to 2,516 generated sample individuals; the 1,647 extraction-eligible rows reconcile to 1,647 DNA extracts; and 1,242 resolved sequence rows reconcile to 1,242 amplicon libraries and 1,242 FASTQ datasets. A site-identity check independently matches each Trip 1 numeric identifier 61–64 to a named catalogue individual by exact coordinates and rejects coordinate or IRI disagreement, accounting for all 36 corresponding source specimens.

Environmental validation accounts for 274 field observations and checks site identity, temperature, pressure and relative humidity. Values enter the curated environmental dataset only when they satisfy the measurement range and identity criteria. The XRF provenance audit reproduced all 13,390 reported Trips 1–4 cells under the maximum-positive rule and all 4,293 reported Trip 5 cells under the last-reported rule, with no mismatched reported cells, and verified that the canonical laboratory input comprises all 725 processed records. An independent chemical-identifier validation evaluated all 93 instrument channels against the pinned ChEBI release 247 and the checksummed PubChem snapshot of 23 July 2026 (Section 2.5); all rows passed the declared rules, with explicit nulls for the aggregate LE channel and unsupported oxides. The pH validation accounts for 1,168 specimen-assay entries and reproduces the recorded admission and exclusion dispositions. It checks specimen identity for every admitted measurement, including field replicate 2 for all 156 Trip 4 measurements.

4.2 Laboratory controls and measurement agreement

The PowerSoil Pro extraction-control profile contained 1,165 ASVs, with 84% of reads assigned to ASVs shared with biological profiles. Its nearest canonical profile was 46Dr1, with Bray–Curtis dissimilarity 0.39. The frozen control adjudication interpreted this profile as contamination of the blank with soil template. Its recorded extraction-blank role and sequence record are retained, with a characterisation-only disposition; it is excluded from the extraction-blank prevalence screen to avoid treating the shared soil signal as reagent contamination. A library-order test compared nearby library numbers with different-site pairs matched by geographic-distance decile and compartment, using 999 library-order permutations. The Trip 1 and Trip 3 Bray–Curtis contrasts were -0.003 ($p = 0.15$) and $+0.006$ ($p = 0.99$), respectively. These contrasts summarize the tested preparation-order association.

To assess agreement between field and laboratory XRF measurements, repeated field sessions were reduced to site medians and paired with the Trip 5 subsurface laboratory observation at 58 sites. Among 12 analytes with positive values in both workflows at more than 10 paired sites, the median Spearman correlation was 0.21 (range -0.40 to 0.48). This diagnostic describes site-level rank agreement among the retained positive pairs.

4.3 Competency questions and SPARQL queries

Competency questions define the intended retrieval tasks (Table 10). The archived query workload covers site, geochemical, molecular-provenance and taxonomic data. The Site 10 XRF query has an archived 46-row result. Source-derived expected result tuples also specify the complete 48-observation genus profile for run `ERR16061083`, 360-row environment–taxon join, 70 site coordinates and six PCR no-template control roles. The release validator compares full result multisets, preserving duplicate rows and requiring exact identifiers. Numeric comparisons use absolute and relative tolerances of 10^{-12}; coordinate comparisons use an absolute tolerance of 10^{-10} degrees. The query comparisons test the links between observations and their specimen, process and site context.

For KG v3.0.1, release validation verifies all 16 input-module hashes and parses the complete 45,707,660-quad export, checking every graph context. The terminology validation checks the term axioms and one definition per annotation property for each of 342 project terms. Public-release checks verify equivalent ontology content in RDF/XML, Turtle, JSON-LD and N-Triples, resolution of the four reference imports, manifest integrity, the live asserted-graph count and the deposited artifacts. Section 3 identifies the validation evidence and its input versions. Structural tests evaluate the scientific modules separately.

Regression fixtures execute the taxonomic and environmental queries in Usage Notes on deliberately confusable inputs. The taxonomy fixture assigns the same genus different abundances in two profiles and adds a species-rank observation; the query returns exactly the two matching genus observations. The visit fixture gives one site two visits at 40 and 20°C; the temperature-filtered query returns only the specimen collected during the 40°C visit. These tests check the implemented joins; release-level validation additionally requires expected result tables for the complete modules.

4.4 MAG catalogue integrity

The integrity check compared the 5,438 refined MAG FASTAs against the representative inventory and the refined-bin CheckM2 report. Every genome's contig count and summed sequence length matched that report. The archive contains 2,485,945 contigs totalling 20,490,828,267 bp. A SHA-256 checksum was computed for each FASTA, and every archive member was reread and its bytes verified against the per-genome manifest. The GTDB R232 classification tables passed their upstream checksum checks and contained exactly the same 5,438 genome identifiers. Library metadata reconciled all 551 read-bearing libraries between the paired-read inventory and the sample-context table.

The gene annotations were checked against the same sequences. Before annotating a genome, we verified its FASTA against the catalogue checksum. We accepted an annotation only when every eggNOG query was a protein that Bakta had predicted for that genome and each query occurred once. All 5,438 genomes passed, and the release step reread every annotation file against its recorded checksum. Bakta predicted 19,186,792 proteins, a median of 3,325 per genome. eggNOG-mapper annotated 84.0% of these proteins and assigned a KO to 44.4%, and every genome has at least one

annotated protein. Across the catalogue, the annotations contain 11,386 distinct KOs, 9,703 Pfam domains and 141 CAZy families.

Each of the 551 coverage tables lists exactly the 5,438 catalogue genomes. All values are finite and non-negative, covered fractions are at most one, read counts are integers, and the reported length of each genome equals its total sequence length in the catalogue. The median library aligned 45.9% of its reads to the catalogue (range 3.9–83.8%). Seven libraries reached the limit of 200 million read pairs. Library F5Dr1, which yielded no assembly, contributed 616 read pairs and had zero mapped reads in the relative-abundance cohort matrix (`EQ-COVERM-REL-ABUND-v1.0.0`). Cohort reconciliation verified exact specimen, site, campaign, compartment and dune-zone joins across the 550 non-zero mapped libraries, the 546 core-transect libraries across 59 primary sites, the 545 read-depth-modeled core libraries, the 442 amplicon-matched libraries across 58 primary sites and the 441 matched read-depth-modeled libraries, together with the strict MIMAG contamination boundary separating the 5,437 ENA MAGs (126 high-quality and 5,311 strict medium-quality representatives) from `48Dr2__self_SemiBin_25` (10.0% contamination) and the separate product provenance of the 20,451,986 Pyrodigal representative proteins (`EQ-PROT-PYRODIGAL-REP-v1.0.0`), the frozen analytical KO matrix (`EQ-KO-LEGACY-ANALYTICAL-v0`), and the frozen analytical taxonomy snapshot (`EQ-TAX-LEGACY-ANALYTICAL-v0`, reconciled 1:1 against canonical GTDB R232 strings across all 5,438 representatives with 1,898 differing strings) associated with the same 5,438 representative genome identifiers.

All 5,438 genomes completed gene calling, profile searches, rule evaluation and continuous-trait prediction. The aggregate tables contain each canonical genome identifier exactly once and retain its GTDB release 232 taxonomy and CheckM2 metrics. The three trait matrices contain 37, 100 and 189 trait categories, respectively, in addition to genome length and continuous predictions. A six-genome smoke test reproduced the official microTrait extraction outputs exactly before the complete cohort was processed. Separate coverage tests verify the missing phosphatase models and prevent unassayed alternatives from being interpreted as absent genes. These checks establish execution and identifier consistency; the biological accuracy of individual trait predictions remains dependent on the underlying models and recovered genome sequence.

4.5 ShEx schema validation

The archived ShEx run used the shapes described in Section 2.7.4. Its tractable modules passed their assigned shapes. The pH module was validated separately over all 712 measurement processes, values and acidity qualities and all 29 sessions, and the laboratory-control module passed its shape and scientific-invariant checks. Four negative fixtures each alter one constraint: missing site coordinates, an invalid value datatype, reversed value-to-quality direction, or an XRF process with neither input nor target. The released shapes reject all four fixtures. The multi-gigabyte taxonomy ABox remains outside this memory-bounded run; it is instead parsed in full by an independent Turtle parser and scanned by a streaming structural validator, which

accounted for 44,528,482 triples in 5,953,690 subject blocks with zero structural violations. This checks full-file syntax and the generator's declared structural relationships. The separate listing validator checks printed Turtle against the modules available for that run. Taxonomy and external reference checks require the corresponding pinned payloads.

4.6 Logical consistency

ELK returned `true` for the tractable scientific-module subset loaded for validation, including the ecosystem, pH and laboratory-control modules. The checked subset excludes the taxonomy ABox and large reference modules. Its execution metadata identify the loaded axioms and import handling; class-hierarchy precomputation was omitted. The ontology uses `xsd:double` literals and constructs outside ELK's supported fragment[57]; the check therefore applies to the constructs processed by that configuration.

Validation of the separate finite-rule implementation uses synthetic examples for long hierarchies, transitive cycles, both project property-chain lengths, restriction and intersection classification, and PCR-control ancestry. The tests compare bounded, class-partitioned execution with an unbounded reference closure and check recovery from an interrupted run. Release-specific checks compare the complete control-role inventory with its source, retrieve consequences of the measurement chains and require the inferred graph to be disjoint from the asserted graph. Completion requires the zero-addition pass specified in Section 2.7.3, followed by the explicit clash diagnostics.

Label validation checks that each labelled project entity has exactly one label. The evidence identifies the input modules, their hashes and validator versions (Section 3).

5 Usage notes

The portal at https://rubalkhali.science/ provides site pages and a SPARQL editor. The endpoint https://rubalkhali.science/sparql accepts SPARQL 1.1 queries. For a local, low-memory example, load the base, site and XRF modules listed in Table 8 and execute the Site 10 query in Listing 8. Compare its 46 returned observations with the distributed expected identifiers and values. Taxonomic queries also require the taxonomy module and its mapping dependencies. The named-run example retrieves the complete genus profile for `ERR16061083`; all-run exports require a locally loaded snapshot and a long-running join. The reproduction instructions accompanying the data provide retrieval, checksum, environment and execution details (Section 3).

5.1 Choosing observations for reuse

Use specimen, collection-visit and assay identifiers when joining observations. The primary repeated-campaign frame comprises sites 1–60; the source identifiers 61–64 refer to four named catalogue locations sampled in Trip 1. Coordinate matching established these aliases, which account for 36 source specimens. Several feature profiles share a source specimen identifier. Analyses should group these profiles by their documented specimen and processing relationships.

The shared cohort and product manifest documents how downstream ecological and proteome analyses subset the resource[19,20]. For 16S amplicon data, excluding the 24 laboratory controls and 10 low-depth biological profiles leaves 1,237 biological profiles, 1,227 primary-transect profiles across sites 1–60, and 1,186 deduplicated primary-transect specimen profiles. For shotgun metagenomics, the 551 read-bearing libraries across 61 site codes (59 primary sites plus Trip 1 named sites 62 and 64) yield 550 non-zero mapped libraries after excluding `F5Dr1`, 546 core-transect libraries across the 59 shotgun-sampled primary sites after excluding the four Trip 1 named-site libraries (`62PRr2`, `64Dr3`, `64PRr2` and `64Sr2`), and 545 core-transect libraries with complete post-QC read-depth covariates after excluding `F45PRr3`. Joining the 546 core-transect shotgun libraries to the 1,186 deduplicated 16S specimen profiles by exact specimen identifier yields 442 matched amplicon–metagenome profiles across 58 primary sites (all primary sites except sites 56 and 59, supporting the 8,860-KO direct comparison with PICRUSt2) and 441 matched profiles with complete post-QC read-depth metadata (`EQ-COVERM-REL-ABUND-v1.0.0`). Competitive archaeal mapping (`EQ-COVERM-ARCHAEA-COMP-v1.0.0`) covers 193 paired shotgun profiles across 68 site-campaign visits and 36 sites against the 5,438 representatives and 301 external Thermoplasmatota genomes across 0%, 10%, 20% and 50% breadth gates, with `MAPQ` $\geq$ 20 evaluated as a separate sensitivity analysis. Users reusing protein annotations or taxonomic labels should select the canonical 19,186,792 Bakta/eggNOG release product (`EQ-PROT-BAKTA-v1.0.0`) and the canonical GTDB R232 taxonomy for new analyses, whereas reproducing frozen companion projections uses the 20,451,986 Pyrodigal representative protein product (`EQ-PROT-PYRODIGAL-REP-v1.0.0`), the frozen analytical KO matrix (`EQ-KO-LEGACY-ANALYTICAL-v0`), or the frozen analytical taxonomy snapshot (`EQ-TAX-LEGACY-ANALYTICAL-v0`) as documented in the shared manifest. The earlier Trips 1–3 pilot dataset (150 CoverM profiles across the 991-genome draft catalogue and its 990-genome KO matrix, `EQ-COVERM-LEGACY-991-v0`) is retained solely for historical reproducibility.

The control registry distinguishes assay, processing stage, extraction group and identity evidence. The Trip 5 candidate screen uses 17 extraction blanks linked to 217 biological profiles and identifies 351 ASVs. The canonical feature table retains these ASVs; the companion sensitivity analysis applies their removal to the linked profiles. Trip 4 has six separately denoised extraction blanks and three batches with unsequenced blanks, leaving complete or partial coverage gaps at 28 sites. Concentration measurements exist in the source records; any concentration-based screen requires links between specimens, extractions and concentration measurements. Four Trips 2–3 standards support genus-recovery evaluation. Trips 4–5 have extraction controls but no sequenced 16S positive-standard profiles.

Archived-soil pH describes specimens measured in $CaCl_2$ under the admission rules of Section 2.5. Coverage varies by campaign and compartment. All Trip 4 assays used archived field replicate 2. Their supported comparison with community profiles is at campaign, site and compartment level. The versioned specimen mapping links measurements to the confirmed specimens as described in Methods.

Field and laboratory XRF observations have separate acquisition provenance and undocumented concentration units. The released values support within-workflow

descriptive comparisons; cross-workflow calibration requires documented units, paired aliquots and reference materials. Site 4 field readings include different substrates, so a site-level summary combines that local heterogeneity with instrument repeatability.

The PMA table contains three site-compartment groups at two campsites. Each group has one parent collection tube, from which three separate slurries supplied treated/untreated pairs. Descriptive paired summaries retain this preparation structure and compare dye treatment under matched dark and light exposure.

5.2 RDF examples and identifier conventions

The examples instantiate the formal patterns with observations from the resource. Table 11 defines the namespaces used in the Turtle examples and SPARQL queries. These declarations must be prepended when parsing an example. The listing validator compares each example with the corresponding modules and pinned reference terms.

Listing 2 represents the Trip 1, Site 1 shallow-subsurface sample. The sampling process targets a site and outputs a collection containing the samples. The sample also links directly to its source site, allowing queries to recover provenance without traversing the sampling event.

Listing 2 Turtle representation of a sampling process.

```
rak:P100001 a rak:0000024 ; # Sampling Process
  sio:000291 rak:1000001 ; # hasTarget
  sio:000229 rak:7000001 . # hasOutput (Collection)

rak:7000001 a sio:001419 ; # Material sample collection
  sio:000059 rak:6000001 . # hasMember

rak:6000001 a rak:0000021 ; # Deep Soil Sample
  sio:000244 rak:1000001 . # isDerivedFrom
```

The same sample is the input of the extraction process in Listing 3. A protocol specification identifies the procedure, and the output DNA extract carries a concentration quality. A separate measurement value quantifies that quality through the reciprocal SIO measurement relations.

Listing 3 Turtle representation of a DNA extraction protocol.

```
rak:P200001 a sio:000994 ; # Scientific Protocol Execution
  sio:000339 rak_kb:DNA_ext_powersoil ; # isSpecifiedBy
  sio:000230 rak:6000001 ; # hasInput (Soil Sample)
  sio:000229 rak:6800001 . # hasOutput (DNA Extract)

rak:6800001 a rak:0000040 ; # DNA Extract
  sio:000008 rak:5900001 . # hasAttribute (Concentration Quality)
```

Taxonomic observations bind a count to both a taxon and its source sequence dataset (Listing 4). Here, sequencing run `ERR16783308` from sample `V10Dr1` has 853,814 reads assigned to Bacteria. Each taxonomic dataset contains one rank and one abundance basis. Count and relative abundance datasets retain the same classifier lineage and source sample context.

Listing 4 Turtle representation of taxonomic abundance formalized from a feature table.

```
rak:P290001 a rak:0000071 ; # Bioinformatic Workflow
  sio:000339 rak:L000012 ; # 16S amplicon processing protocol
  sio:000230 rak:7791009 ; # hasInput (FASTQ Dataset, ERR16783308)
  sio:000229 rak:7740001 . # hasOutput (Taxonomic Dataset)

rak:7740001 a rak:0000074 ; # taxon absolute abundance dataset
  sio:000059 rak:V00000001 . # hasMember (Observation Value)

rak:V00000001 a rak:0000076 ; # Absolute abundance value
  sio:000215 rak:Q00000001 ; # isMeasurementValueOf
  rak:2000025 "Domain: Bacteria" ; # classifier lineage
  rak:2000026 "853814.0"^^xsd:double . # absolute read count

rak:Q00000001 a rak:0000078, pato:0000070 ; # Absolute Abundance Quality
  sio:000011 ncbitaxon:2, rak:7791009 ; # taxon and FASTQ bearers
  sio:000216 rak:V00000001 . # hasMeasurementValue
```

Laboratory XRF specializes the measurement pattern by linking the process to the archived sample from which it was measured. Listing 5 shows the aluminium observation for `V1Sr1`. Its numerical value and quality have reciprocal links. The value retains the instrument reported magnitude with undocumented concentration units. Field sessions instead identify a target site through `sio:000291`.

Listing 5 Turtle representation of an XRF geochemistry measurement.

```
rak:P300072 a rak:0000025 ; # XRF Measuring Process
  sio:000230 rak:6001732 ; # hasInput (sample V1Sr1)
  sio:000339 rak:L_Surface_Lab ; # laboratory protocol
  sio:000229 rak:4302121 . # hasOutput

rak:4302121 a rak:0000507 ; # Al concentration value
  sio:000215 rak:5302121 ; # isMeasurementValueOf
  rak:2000012 "0.78"^^xsd:double . # Concentration

rak:5302121 a rak:0000107 ; # Al concentration quality
  sio:000011 rak:6001732 ; # isAttributeOf
  sio:000216 rak:4302121 . # hasMeasurementValue
```

Shared inputs and outputs connect extraction, library preparation and sequencing (Listing 6). The final sequencing dataset links to its ENA accession. A SPARQL traversal can therefore recover the recorded provenance from sample to sequencing run.

Listing 6 Turtle representation of the sequencing workflow from soil sample T1Dr1 to ENA run ERR16061083.

```
rak:P200923 a sio:000994 ; # DNA Extraction Process
  sio:000339 rak_kb:DNA_ext_powersoil ; # isSpecifiedBy
  sio:000230 rak:6000577 ; # hasInput (T1Dr1 Soil Sample)
  sio:000229 rak:6800540 . # hasOutput (DNA Extract)

rak:P250001 a rak:0000065 ; # Library Preparation Process
  sio:000339 rak:L000010 ; # Illumina 16S Prep Guide
  sio:000230 rak:6800540 ; # hasInput (DNA Extract)
  sio:000229 rak:6690001 . # hasOutput (Amplicon Library)

rak:P280001 a rak:0000066 ; # Sequencing Process
  sio:000230 rak:6690001 ; # hasInput (Amplicon Library)
  sio:000229 rak:7790001 . # hasOutput (FASTQ Dataset)
```

```
rak:7790001 a rak:0000063 ; # FASTQ Dataset
  rdfs:seeAlso <https://identifiers.org/insdc.run/ERR16061083> .
```

Quality control observations retain the read direction through separate forward and reverse metric classes. Listing 7 gives the forward-read GC content for run ERR16062319: 58.0%. The quality belongs to the same FASTQ dataset consumed by the quality control process.

Listing 7 Turtle representation of forward-read GC content for ENA run ERR16062319.

```
rak:P350001 a rak:0000200 ; # Sequencing QC Process
  sio:000230 rak:7790263 ; # hasInput (FASTQ Dataset)
  sio:000229 rak:4450001 . # hasOutput (QC Value)

rak:4450001 a rak:0000217 ; # Forward GC Content Value
  sio:000215 rak:5550001 ; # isMeasurementValueOf
  sio:000221 uo:0000187 ; # Percent
  rak:2000074 "58.0"^^xsd:double . # has GC content

rak:5550001 a rak:0000205 ; # Forward GC Content Quality
  sio:000011 rak:7790263 ; # isAttributeOf (FASTQ Dataset)
  sio:000216 rak:4450001 . # hasMeasurementValue
```

5.3 Retrieving measurements and community profiles

Listing 8 retrieves XRF measurements for a specific site, requiring traversal of the measurement process, quality, and value entities.

Listing 8 SPARQL query to retrieve field XRF measurements for Site 10. The query starts from the field-XRF measuring processes whose target is the site and walks to the individual analyte concentrations.

```
PREFIX rak: <https://rubalkhali.science/kb/RAK_>
PREFIX sio: <http://semanticscience.org/resource/SIO_>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>

SELECT ?processLabel ?siteLabel ?analyte ?concentration
WHERE {
  ?process a rak:0000025 ;
           rdfs:label ?processLabel ;
           sio:000291 ?site ;
           sio:000229 ?value .
  FILTER(CONTAINS(LCASE(STR(?processLabel)), "field xrf"))
  ?site rdfs:label ?siteLabel .
  FILTER(?siteLabel = "Site 10")
  ?value sio:000215 ?quality ;
         rak:2000012 ?concentration .
  ?quality a ?qualityClass .
  ?qualityClass rdfs:label ?analyte .
  FILTER(?qualityClass != rak:0000032)
}
ORDER BY ?processLabel ?analyte
```

The query selects field-XRF processes targeting Site 10, then retrieves each output and its analyte quality. It returns one row per process/analyte combination, excluding the aggregate Light Elements channel. The instrument-reported values have undocumented units.

The query returned 46 analyte-value rows for Site 10: 24 from Test 5847 and 22 from Test 5848. Validation compares the returned identifiers and values with the expected observations from the site and XRF data.

Listing 9 retrieves the complete genus profile for run `ERR16061083`, joining across bioinformatic workflows, FASTQ datasets, and taxon observations.

Listing 9 Complete genus-level profile for run ERR16061083, ordered by descending read count and lineage.

```
PREFIX rak: <https://rubalkhali.science/kb/RAK_>
PREFIX sio: <http://semanticscience.org/resource/SIO_>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
PREFIX xsd: <http://www.w3.org/2001/XMLSchema#>

SELECT ?proc ?fastq ?runLabel ?lineage ?taxon ?taxonLabel
       ?count ?relativeAbundance
WHERE {
  VALUES ?runLabel { "FASTQ dataset for ERR16061083" }
  ?proc a rak:0000071 ;
        sio:000230 ?fastq ;
        sio:000229 ?dsAbs, ?dsRel .
  ?fastq rdfs:label ?runLabel .
  ?dsAbs a rak:0000074 ; rdfs:label ?absLabel ;
         sio:000059 ?valAbs .
  ?dsRel a rak:0000075 ; rdfs:label ?relLabel ;
         sio:000059 ?valRel .
  FILTER(STRENDS(STR(?absLabel), "(Genus)"))
  FILTER(STRENDS(STR(?relLabel), "(Genus)"))
  ?valAbs rak:2000026 ?count ;
          rak:2000025 ?lineage ;
          sio:000215 ?qualAbs .
  ?valRel rak:2000020 ?relativeAbundance ;
          rak:2000025 ?lineage ;
          sio:000215 ?qualRel .
  ?qualAbs sio:000011 ?taxon, ?fastq .
  ?qualRel sio:000011 ?taxon, ?fastq .
  ?taxon rdfs:label ?taxonLabel .
  FILTER(?taxon != ?fastq)
}
ORDER BY ?runLabel DESC(xsd:double(?count)) ?lineage ?proc
```

The query binds both count and relative-abundance values to the same workflow, source FASTQ dataset, taxon and classifier lineage. Dataset labels encode rank through the suffix `(Genus)`. The query uses that emitted convention explicitly and returns all 48 observations for the named run. The source-derived expected table contains the complete profile, including each classifier lineage. The all-run query retrieves profiles from a locally loaded snapshot. That unrestricted export is a long-running operation. The environmental example additionally binds field temperature and specimen collection to the same site visit.

Listing 10 SPARQL query to retrieve Pseudomonas relative abundance at sites with high temperature.

```
PREFIX rak: <https://rubalkhali.science/kb/RAK_>
PREFIX sio: <http://semanticscience.org/resource/SIO_>
PREFIX pato: <http://purl.obolibrary.org/obo/PATO_>
PREFIX ncbitaxon: <http://purl.obolibrary.org/obo/NCBITaxon_>
PREFIX rdfs: <http://www.w3.org/2000/01/rdf-schema#>
SELECT ?visit ?visitLabel ?site ?siteLabel ?sample ?fastq
```

```
      ?temperatureProcess ?tempVal ?temp ?lineage ?relAbundance
WHERE {
  ?visit a rak:0000003 ; rdfs:label ?visitLabel ;
         sio:000291 ?site .
  ?site rdfs:label ?siteLabel .
  ?temperatureProcess sio:000068 ?visit ;
                      sio:000229 ?tempVal .
  ?tempVal rak:2000003 ?temp ; sio:000215 ?tempQual .
  ?tempQual a pato:0000146 ; sio:000011 ?site .
  FILTER(?temp > 35.0)
  ?sampling a rak:0000024 ; sio:000068 ?visit ;
            sio:000291 ?site ; sio:000229 ?collection .
  ?collection sio:000059 ?sample .
  ?ext sio:000230 ?sample ; sio:000229 ?dna .
  ?libPrep a rak:0000065 ; sio:000230 ?dna ;
           sio:000229 ?lib .
  ?seq a rak:0000066 ; sio:000230 ?lib ;
       sio:000229 ?fastq .
  ?bioinf a rak:0000071 ; sio:000230 ?fastq ;
          sio:000229 ?dsRel .
  ?dsRel a rak:0000075 ; rdfs:label ?datasetLabel ;
         sio:000059 ?valRel .
  FILTER(STRENDS(STR(?datasetLabel), "(Genus)"))
  ?valRel rak:2000020 ?relAbundance ;
          rak:2000025 ?lineage ; sio:000215 ?qualRel .
  ?qualRel sio:000011 ncbitaxon:286, ?fastq .
}
ORDER BY ?visit ?sample ?fastq ?tempVal ?lineage
```

Each result retains the collection visit, specimen, sequence dataset, temperature observation and genus lineage. A site's visits are joined through their process provenance. This provides an environmental temperature associated with the collection visit of each returned microbial profile.

5.4 Versioning and supported operations

Use the versioned modules and their reference ontologies to reproduce queries. Updating an external ontology requires revalidating identifiers and mappings before regenerating the affected modules and query results.

KG v3.0.1 serves the asserted graph at https://rubalkhali.science/graph/asserted/3.0.1. This graph is the default dataset for portal and SPARQL queries. Explicit `FROM`, `FROM NAMED` and SPARQL protocol dataset parameters select the dataset for an individual request. The service evaluates stored assertions directly and rejects unavailable query modes.

The bundled finite-rule implementation supports the separate reasoning workflow described in Section 2.7.3. Publication of its output requires completion and validation of that workflow. Changes to released artifacts require a new semantic version.

A local copy of the deposited graph provides fixed inputs independently of the portal's deployed version. The rebuild kit verifies and loads the 16 distributed modules; generating the scientific modules uses the source data and code described in Section 3. Analyses requiring additional OWL consequences need a reasoner and validation workload suited to those constructs.

6 Data availability

The project uses ENA umbrella study PRJEB104209; its submitted amplicon run records are associated with PRJEB106069. ENA study PRJEB125153 (secondary accession ERP204894) holds the shotgun data: 554 read runs from 551 soil BioSamples, 549 primary metagenome assemblies, the 12,434 refined bins as binned metagenomes, and 5,437 of the 5,438 representatives as MAGs with GCA accessions. ENA admits a genome as a MAG only below 10% contamination, so 48Dr2__self_SemiBin_25 is deposited as a binned metagenome. Both ENA studies have assigned accessions, with public release scheduled for 25 November 2026. The data repository lists the BioSample, run, assembly and genome accessions of every shotgun library, bin and MAG. The complete refined MAG sequence archive, its functional annotation tables, the coverage profiles and the accession tables are prepared and checksum-verified in the release candidate; Zenodo deposition and DOI assignment (DATA-BULK-ARCHIVE-DOI) are pending verification[58], after which the completed archive link will be published at https://rubalkhali.science/. The source and derived-data package is maintained at https://github.com/bio-ontology-research-group/empty-quarter-ecology-data. Its manifests distinguish included files, separately retrieved bulk payloads and unresolved upstream provenance. KG v3.0.1, its original modules and rebuild kit are archived at https://doi.org/10.5281/zenodo.23134168[55]. Project-authored RDF, ontology terms, definitions and dataset metadata in this release are licensed under CC BY 4.0. Imported vocabularies retain their respective licences, listed in the release's third-party attribution file; project software is MIT-licensed. Raw reads and other separately licensed materials retain their own access and licensing conditions. The broader analytical tables and specimen and assay provenance are versioned in the reproducibility repository; the KG deposit contains the files enumerated in its manifest.

7 Code availability

The archived KG v3.0.1 release includes the rebuild kit and its version-specific manifest[55]. The web application and ontology-development sources are maintained at https://github.com/bio-ontology-research-group/empty-quarter, and the data-processing sources at https://github.com/bio-ontology-research-group/empty-quarter-ecology-data, which contains the manuscript sources, RDF generators, metadata and Nextflow workflow. RDF generation uses Groovy scripts with OWL API v5.1.20[49]; the ELK validation wrapper uses OWL API v4.5.26. The repository supplies execution profiles, workflow tests, a pinned Conda recipe and a hash-locked Python environment. Project software is MIT-licensed. The release manifest identifies the archived rebuild files; the accompanying repositories document their execution environments.

Funding

Susana Martinez Arbas was supported by the Fonds National de la Recherche (FNR), Luxembourg (Project Ref. 18047750). No specific funding was received for this project.

Competing interests

The authors declare no competing interests.

References

[1] Mandaville, J.P.: Plant life in the rub' al-khali (the empty quarter), south-central arabia. Proceedings of the Royal Society of Edinburgh. Section B. Biological Sciences **89**, 147–157 (1986) https://doi.org/10.1017/s0269727000008988

[2] Edgell, H.S.: Arabian Deserts: Nature, Origin, and Evolution. Springer, ??? (2006)

[3] Nelli, N., Francis, D., Sow, M., Fonseca, R., Alkatheeri, A., Bosc, E., Bergametti, G.: The wind-blown sand experiment in the empty quarter desert: Roughness length and saltation characteristics. Earth and Space Science **11**(6), 2024–003512 (2024) https://doi.org/10.1029/2024EA003512

[4] Makhalanyane, T.P., Valverde, A., Gunnigle, E., Frossard, A., Ramond, J.-B., Cowan, D.A.: Microbial ecology of hot desert edaphic systems. FEMS Microbiology Reviews **39**(2), 203–221 (2015) https://doi.org/10.1093/femsre/fuu011

[5] Yasir, M., Azhar, E.I., Khan, I., Bibi, F., Baabdullah, R., Al-Zahrani, I.A., Al-Ghamdi, A.K.: Composition of soil microbiome along elevation gradients in southwestern highlands of Saudi Arabia. BMC Microbiology **15**(1), 65 (2015) https://doi.org/10.1186/s12866-015-0398-4

[6] Aslam, S.N., Dumbrell, A.J., Sabir, J.S., Mutwakil, M.H.Z., Baeshen, M.M.N., Abo-Aba, S.E.M., Clark, D.R., Yates, S.A., Baeshen, N.A., Underwood, G.J.C., McGenity, T.J.: Soil compartment is a major determinant of the impact of simulated rainfall on desert microbiota. Environmental Microbiology **18**(12), 5048–5062 (2016) https://doi.org/10.1111/1462-2920.13474

[7] Khan, M.A., Khan, S.T.: Microbial communities and their predictive functional profiles in the arid soil of Saudi Arabia. SOIL **6**, 513–521 (2020) https://doi.org/10.5194/soil-6-513-2020

[8] Eida, A.A., Ziegler, M., Lafi, F.F., Michell, C.T., Voolstra, C.R., Hirt, H., Saad, M.M.: Desert plant bacteria reveal host influence and beneficial plant growth properties. PLOS ONE **13**(12), 0208223 (2018) https://doi.org/10.1371/journal.pone.0208223

[9] Pesant, S., Not, F., Picheral, M., Kandels-Lewis, S., Le Bescot, N., Gorsky, G., Iudicone, D., Karsenti, E., Speich, S., Troublé, R., Dimier, C., Searson, S., Tara Oceans Consortium Coordinators: Open science resources for the discovery and analysis of Tara Oceans data. Scientific Data **2**, 150023 (2015) https://doi.org/

10.1038/sdata.2015.23

[10] Yilmaz, P., Kottmann, R., Field, D., Knight, R., Cole, J.R., Amaral-Zettler, L., Gilbert, J.A., *et al.*: Minimum information about a marker gene sequence (MIMARKS) and minimum information about any (x) sequence (MIxS) specifications. Nature Biotechnology **29**(5), 415–420 (2011) https://doi.org/10.1038/nbt.1823

[11] National Microbiome Data Collaborative: National Microbiome Data Collaborative: Data Standards. https://microbiomedata.org/data-standards/. Accessed 23 July 2026 (2026)

[12] Santangelo, B.E., Hegde, H., Caufield, J.H., Reese, J., Kliegr, T., Hunter, L.E., Lozupone, C.A., Mungall, C.J., Joachimiak, M.P.: KG-Microbe – building modular and scalable knowledge graphs for microbiome and microbial sciences. GigaScience, 077 (2026) https://doi.org/10.1093/gigascience/giag077

[13] Ma, C., Liu, S., Won, S., Koslicki, D.: MetagenomicKG: a knowledge graph for metagenomic applications. Bioinformatics **42**(7), 421 (2026) https://doi.org/10.1093/bioinformatics/btag421

[14] Dumontier, M., Baker, C.J., Baran, J., Callahan, A., Chepelev, L., Cruz-Toledo, J., Del Rio, N.R., Duck, G., Furlong, L.I., Keath, N., *et al.*: The Semanticscience Integrated Ontology (SIO) for biomedical research and knowledge discovery. Journal of Biomedical Semantics **5**(1), 14 (2014) https://doi.org/10.1186/2041-1480-5-14

[15] Degtyarenko, K., Matos, P., Ennis, M., Hastings, J., Zbinden, M., McNaught, A., Alcántara, R., Darsow, M., Guedj, M., Ashburner, M.: ChEBI: a database and ontology for chemical entities of biological interest. Nucleic Acids Research **36**(suppl_1), 344–350 (2008) https://doi.org/10.1093/nar/gkm791

[16] Buttigieg, P.L., Morrison, N., Smith, B., Mungall, C.J., Lewis, S.E.: The environment ontology: contextualising biological and biomedical entities. Journal of Biomedical Semantics **4**(1), 43 (2013)

[17] Schoch, C.L., Ciufo, S., Domrachev, M., Hotton, C.L., Kannan, S., Khovanskaya, R., Leipe, D., Mcveigh, R., O'Neill, K., Robbertse, B., *et al.*: NCBI Taxonomy: a comprehensive update on curation, resources and tools. Database **2020**, 062 (2020) https://doi.org/10.1093/database/baaa062

[18] Kazakov, Y., Krötzsch, M., Simančík, F.: The incredible ELK: From polynomial procedures to efficient reasoning with $\mathcal{EL}$ ontologies. Journal of Automated Reasoning **53**(1), 1–61 (2014) https://doi.org/10.1007/s10817-013-9296-3

[19] Tawfiq, R., Abdelhakim, M., Alajel, S.M., Alarawi, M., Aldakhil, H., Derouiche,

A., Drautz-Moses, D.I., Dumontier, M., Grünberg, R., Lopez-Velazquez, A., Martinez Arbas, S., Niu, K., Sivakumar, K., Wang, T., Zhao, X., Zubair, J., Rueping, M., Hoehndorf, R.: Microbial biogeography and functional differentiation across the Rub' al-Khali. Companion ecology manuscript in preparation (2026)

[20] Lopez Velazquez, A., Abdelhakim, M., Alajel, S.M., Alarawi, M., Aldakhil, H., Derouiche, A., Drautz-Moses, D.I., Dumontier, M., Grünberg, R., Kulmanov, M., Martinez Arbas, S., Niu, K., Sivakumar, K., Tawfiq, R., Wang, T., Zhao, X., Zubair, J.K., Rueping, M., Hoehndorf, R.: Widespread desert archaea harbour an export-enriched unknown protein repertoire. Companion proteome manuscript in preparation (2026)

[21] Klindworth, A., Pruesse, E., Schweer, T., Peplies, J., Quast, C., Horn, M., Glöckner, F.O.: Evaluation of general 16s ribosomal RNA gene PCR primers for classical and next-generation sequencing-based diversity studies. Nucleic Acids Research **41**(1), 1 (2013) https://doi.org/10.1093/nar/gks808

[22] Illumina: 16s metagenomic sequencing library preparation: Preparing 16s ribosomal rna gene amplicons for the illumina miseq system (2013). Available at https://support.illumina.com/documents/documentation/chemistry_documentation/16s/16s-metagenomic-library-prep-guide-15044223-b.pdf; accessed 27 August 2026

[23] Fierer, N., Leung, P.M., Lappan, R., Eisenhofer, R., Ricci, F., Holland, S.I., *et al.*: Guidelines for preventing and reporting contamination in low-biomass microbiome studies. Nature Microbiology **10**(7), 1570–1580 (2025) https://doi.org/10.1038/s41564-025-02035-2

[24] Carini, P., Marsden, P.J., Leff, J.W., Morgan, E.E., Strickland, M.S., Fierer, N.: Relic dna is abundant in soil and obscures estimates of soil microbial diversity. Nature microbiology **2**(3), 16242 (2016)

[25] Ewels, P.A., Peltzer, A., Fillinger, S., Patel, H., Alneberg, J., Wilm, A., Garcia, M.U., Di Tommaso, P., Nahnsen, S.: The nf-core framework for community-curated bioinformatics pipelines. Nature Biotechnology **38**(3), 276–278 (2020) https://doi.org/10.1038/s41587-020-0439-x

[26] Straub, D., Blackwell, N., Langarica-Fuentes, A., Peltzer, A., Nahnsen, S., Kleindienst, S.: Interpretations of environmental microbial community studies are biased by the selected 16s rrna (gene) amplicon sequencing pipeline. Frontiers in Microbiology **11**, 550420 (2020)

[27] Callahan, B.J., McMurdie, P.J., Rosen, M.J., Han, A.W., Johnson, A.J.A., Holmes, S.P.: DADA2: high-resolution sample inference from Illumina amplicon data. Nature Methods **13**(7), 581–583 (2016) https://doi.org/10.1038/nmeth.3869

[28] Quast, C., Pruesse, E., Yilmaz, P., Gerken, J., Schweer, T., Yarza, P., Peplies, J., Glöckner, F.O.: The SILVA ribosomal RNA gene database project: improved data processing and web-based tools. Nucleic Acids Research **41**(D1), 590–596 (2013) https://doi.org/10.1093/nar/gks1219

[29] Douglas, G.M., Maffei, V.J., Zaneveld, J.R., Yurgel, S.N., Brown, J.R., Taylor, C.M., Huttenhower, C., Langille, M.G.I.: PICRUSt2 for prediction of metagenome functions. Nature Biotechnology **38**(6), 685–688 (2020) https://doi.org/10.1038/s41587-020-0548-6

[30] Kim, S., Chen, J., Cheng, T., Gindulyte, A., He, J., He, S., Li, Q., Shoemaker, B.A., Thiessen, P.A., Yu, B., *et al.*: PubChem 2023 update. Nucleic Acids Research **51**(D1), 1373–1380 (2023)

[31] Open-Meteo: Open-Meteo: free weather API. https://open-meteo.com/ (2024)

[32] Li, D., Liu, C.-M., Luo, R., Sadakane, K., Lam, T.-W.: MEGAHIT: an ultra-fast single-node solution for large and complex metagenomics assembly via succinct de Bruijn graph. Bioinformatics **31**(10), 1674–1676 (2015)

[33] Pan, S., Zhu, C., Zhao, X.-M., Coelho, L.P.: SemiBin2: self-supervised contrastive learning leads to better MAGs for short- and long-read sequencing. Bioinformatics **39**(Supplement_1), 21–29 (2023)

[34] Sieber, C.M.K., Probst, A.J., Sharrar, A., Thomas, B.C., Hess, M., Tringe, S.G., Banfield, J.F.: Recovery of genomes from metagenomes via a dereplication, aggregation and scoring strategy. Nature Microbiology **3**(7), 836–843 (2018) https://doi.org/10.1038/s41564-018-0171-1

[35] Chklovski, A., Parks, D.H., Woodcroft, B.J., Tyson, G.W.: CheckM2: a rapid, scalable and accurate tool for assessing microbial genome quality using machine learning. Nature Methods **20**(8), 1203–1212 (2023) https://doi.org/10.1038/s41592-023-01940-w

[36] Olm, M.R., Brown, C.T., Brooks, B., Banfield, J.F.: drep: a tool for fast and accurate genomic comparisons that enables improved genome recovery from metagenomes through de-replication. The ISME Journal **11**(12), 2864–2868 (2017)

[37] Parks, D.H., Chuvochina, M., Rinke, C., Mussig, A.J., Chaumeil, P.-A., Hugenholtz, P.: GTDB: an ongoing census of bacterial and archaeal diversity through a phylogenetically consistent, rank normalized and complete genome-based taxonomy. Nucleic Acids Research **50**(D1), 785–794 (2022) https://doi.org/10.1093/nar/gkab776

[38] Chan, P.P., Lin, B.Y., Mak, A.J., Lowe, T.M.: tRNAscan-SE 2.0: improved detection and functional classification of transfer RNA genes. Nucleic Acids Research

49(16), 9077–9096 (2021) https://doi.org/10.1093/nar/gkab688

[39] Bowers, R.M., Kyrpides, N.C., Stepanauskas, R., Harmon-Smith, M., Doud, D., Reddy, T.B.K., Schulz, F., Jarett, J., Rivers, A.R., Eloe-Fadrosh, E.A., Tringe, S.G., *et al.*: Minimum information about a single amplified genome (MISAG) and a metagenome-assembled genome (MIMAG) of bacteria and archaea. Nature Biotechnology **35**(8), 725–731 (2017) https://doi.org/10.1038/nbt.3893

[40] Schwengers, O., Jelonek, L., Dieckmann, M.A., Beyvers, S., Blom, J., Goesmann, A.: Bakta: rapid and standardized annotation of bacterial genomes via alignment-free sequence identification. Microbial Genomics **7**(11), 000685 (2021) https://doi.org/10.1099/mgen.0.000685

[41] Cantalapiedra, C.P., Hernández-Plaza, A., Letunic, I., Bork, P., Huerta-Cepas, J.: eggNOG-mapper v2: functional annotation, orthology assignments, and domain prediction at the metagenomic scale. Molecular Biology and Evolution **38**(12), 5825–5829 (2021) https://doi.org/10.1093/molbev/msab293

[42] Langmead, B., Salzberg, S.L.: Fast gapped-read alignment with Bowtie 2. Nature Methods **9**(4), 357–359 (2012) https://doi.org/10.1038/nmeth.1923

[43] Danecek, P., Bonfield, J.K., Liddle, J., Marshall, J., Ohan, V., Pollard, M.O., Whitwham, A., Keane, T., McCarthy, S.A., Davies, R.M., Li, H.: Twelve years of SAMtools and BCFtools. GigaScience **10**(2), 008 (2021) https://doi.org/10.1093/gigascience/giab008

[44] Aroney, S.T.N., Newell, R.J.P., Nissen, J.N., Camargo, A.P., Tyson, G.W., Woodcroft, B.J.: CoverM: read alignment statistics for metagenomics. Bioinformatics **41**(4), 147 (2025) https://doi.org/10.1093/bioinformatics/btaf147

[45] W3C OWL Working Group: OWL 2 Web Ontology Language Document Overview (Second Edition). W3c recommendation, W3C (December 2012). https://www.w3.org/TR/owl2-overview/

[46] Gkoutos, G.V., Schofield, P.N., Hoehndorf, R.: The Units Ontology: a tool for integrating units of measurement in science. Database **2012**, 033 (2012)

[47] Balvočiūtė, M., Huson, D.H.: Silva, rdp, greengenes, ncbi and ott—how do these taxonomies compare? BMC genomics **18**(Suppl 2), 114 (2017)

[48] Baader, F., Calvanese, D., McGuinness, D.L., Nardi, D., Patel-Schneider, P.F.: The Description Logic Handbook: Theory, Implementation, and Applications, 2nd edn. Cambridge University Press, Cambridge (2007). https://doi.org/10.1017/CBO9780511711787

[49] Horridge, M., Bechhofer, S.: The OWL API: a Java API for OWL ontologies. Semantic Web **2**(1), 11–21 (2011)

[50] Cuenca Grau, B., Horrocks, I., Kazakov, Y., Sattler, U.: Modular reuse of ontologies: theory and practice. Journal of Artificial Intelligence Research **31**, 273–318 (2008) https://doi.org/10.1613/jair.2375

[51] OpenLink Software: OpenLink Virtuoso Universal Server. https://virtuoso.openlinksw.com/ (2024)

[52] Prud'hommeaux, E., Labra Gayo, J.E., Solbrig, H.: Shape Expressions: an RDF validation and transformation language. In: Proceedings of the 10th International Conference on Semantic Systems (SEM '14), pp. 32–40. ACM, Leipzig, Germany (2014). https://doi.org/10.1145/2660517.2660523

[53] Karaoz, U., Brodie, E.L.: microTrait: A toolset for a trait-based representation of microbial genomes. Frontiers in Bioinformatics **2**, 918853 (2022) https://doi.org/10.3389/fbinf.2022.918853

[54] Di Tommaso, P., Chatzou, M., Floden, E.W., Barja, P.P., Palumbo, E., Notredame, C.: Nextflow enables reproducible computational workflows. Nature biotechnology **35**(4), 316–319 (2017)

[55] Tawfiq, R., et al.: The microbial diversity of Rub' al-Khali (Empty Quarter). Zenodo (2026). https://doi.org/10.5281/zenodo.23134168 . https://doi.org/10.5281/zenodo.23134168

[56] Wilkinson, M.D., Dumontier, M., Aalbersberg, I.J., Appleton, G., Axton, M., Baak, A., Blomberg, N., Boiten, J.-W., Silva Santos, L.B., Bourne, P.E., *et al.*: The FAIR Guiding Principles for scientific data management and stewardship. Scientific Data **3**(1), 160018 (2016)

[57] Motik, B., Cuenca Grau, B., Horrocks, I., Wu, Z., Fokoue, A., Lutz, C.: OWL 2 Web Ontology Language Profiles (Second Edition). W3c recommendation, W3C (December 2012). https://www.w3.org/TR/owl2-profiles/

[58] Tawfiq, R., et al.: Metagenome-assembled genome catalogue of Rub' al-Khali (Empty Quarter) desert soils. Dataset archive in preparation; Zenodo deposition and DOI (`DATA-BULK-ARCHIVE-DOI`) pending (2026)

Table 9 Principal products and their representation. Assigned product identifiers denote manuscript and reproducibility snapshot versions distinguishing the canonical resource from companion analytical products and historical pilots. Counts refer to the indicated entity for each product. Archive locations and access conditions are given in Data Availability.

Product	Entity count	Format and graph coverage
Specimen catalogue	2,516 non-control entries	Tables and RDF specimen/process links
Amplicon profiles	1,271 profiles (1,186 deduplicated primary-transect specimens)	Count tables; taxonomic observations and provenance in RDF
Shotgun reads	551 libraries; 554 runs	ENA reads and accession tables; sequence metadata in RDF
Primary assemblies	549 assemblies	ENA sequences and provenance tables
Representative genomes	5,438 genomes (126 HQ, 5,311 strict MQ, 1 at 10.0% contamination)	FASTA and quality/taxonomy tables; separate from the 16 RDF modules
Bakta/eggNOG annotations	19,186,792 proteins (16,119,759 annotated; 8,516,493 KO across 11,386 KOs, 9,703 Pfam, 141 CAZy)	Gene and genome feature tables, tool versions and checksums (`EQ-PROT-BAKTA-v1.0.0`)
Companion Pyrodigal proteins	20,451,986 representative proteins (20,379,930 bin-mapped)	Protein sequences and reconciliation summaries (`EQ-PROT-PYRODIGAL-REP-v1.0.0`)
microTrait genome traits	5,438 genomes; 37, 100, 189 categories	Compressed model, rule, trait and continuous-prediction tables (`EQ-MICROTRAIT-REP-v1.0.0`); separate from KG v3.0.1
Coverage and cohort matrices	551 depth/breadth; 550, 546, 545, 442 and 441 cohorts; 193 archaeal; 150 pilot	Snapshot CoverM products (`EQ-COVERM-REFINED-v1.0.0`, `REL-ABUND-v1.0.0`, `ARCHAEA-COMP-v1.0.0`, `LEGACY-991-v0`)
Soil pH	712 admitted measurements	Tables and RDF, with excluded-row dispositions and Trip 4 specimen links (`EQ-PH-SHARED-v1.0.1`)
XRF	725 laboratory records; 71 field sessions	Tables and RDF; acquisition workflows kept separate
Cohort/product manifest	5,438 genomes; 551 shotgun and 1,186 amplicon cohort flags	Machine-readable provenance manifest, per-library accounting and MIMAG boundary table
Asserted knowledge graph	16 modules (45,707,660 triples)	RDF/XML, Turtle and combined compressed N-Quads

Table 10 Competency questions defining the retrieval workload.

ID	Competency question
CQ1	What are the XRF analyte concentrations measured at a given sampling site?
CQ2	Which sampling sites are recorded, and what are their GPS coordinates?
CQ3	Which biome and environmental features are assigned to each sampling site?
CQ4	What are the instrument-reported Light Elements values across the XRF measurements?
CQ5	Which DNA extraction records (extract, concentration, protocol) exist for a given sample?
CQ6	What is the full provenance chain from a soil sample to its sequencing run?
CQ7	What are the complete genus-level read counts and relative read abundance for run `ERR16061083`?

Table 11 Prefix declarations used in all Turtle and SPARQL listings.

Prefix	Namespace
`rak:`	`https://rubalkhali.science/kb/RAK_`
`rak_kb:`	`https://rubalkhali.science/kb/`
`sio:`	`http://semanticscience.org/resource/SIO_`
`owl:`	`http://www.w3.org/2002/07/owl#`
`pato:`	`http://purl.obolibrary.org/obo/PATO_`
`uo:`	`http://purl.obolibrary.org/obo/UO_`
`envo:`	`http://purl.obolibrary.org/obo/ENVO_`
`ncbitaxon:`	`http://purl.obolibrary.org/obo/NCBITaxon_`
`obo:`	`http://purl.obolibrary.org/obo/`
`rdf:`	`http://www.w3.org/1999/02/22-rdf-syntax-ns#`
`rdfs:`	`http://www.w3.org/2000/01/rdf-schema#`
`xsd:`	`http://www.w3.org/2001/XMLSchema#`
`geo:`	`http://www.opengis.net/ont/geosparql#`
`dcterms:`	`http://purl.org/dc/terms/`